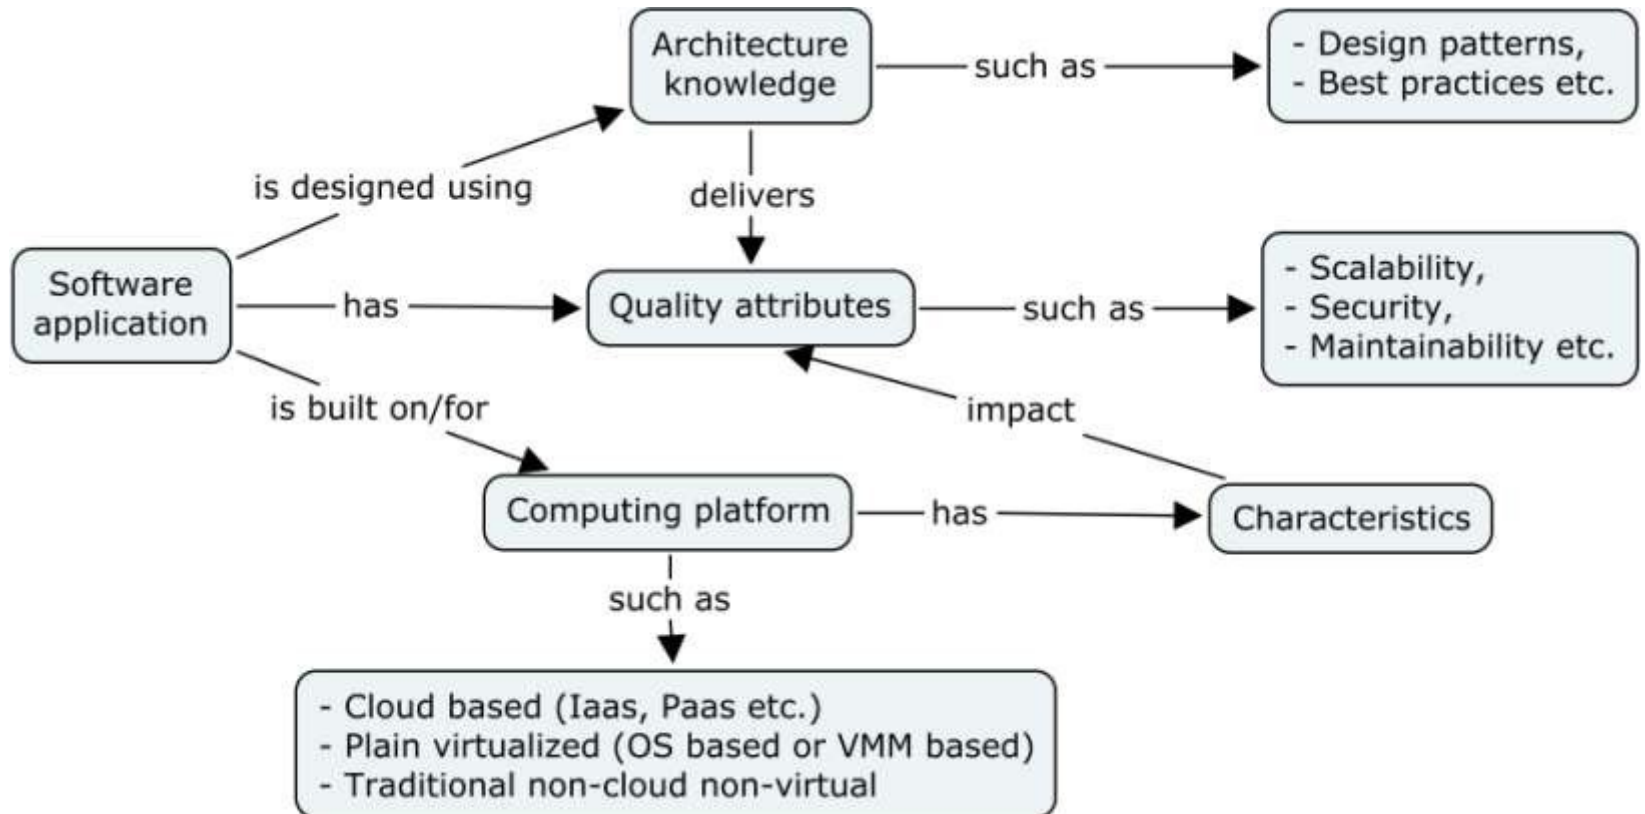
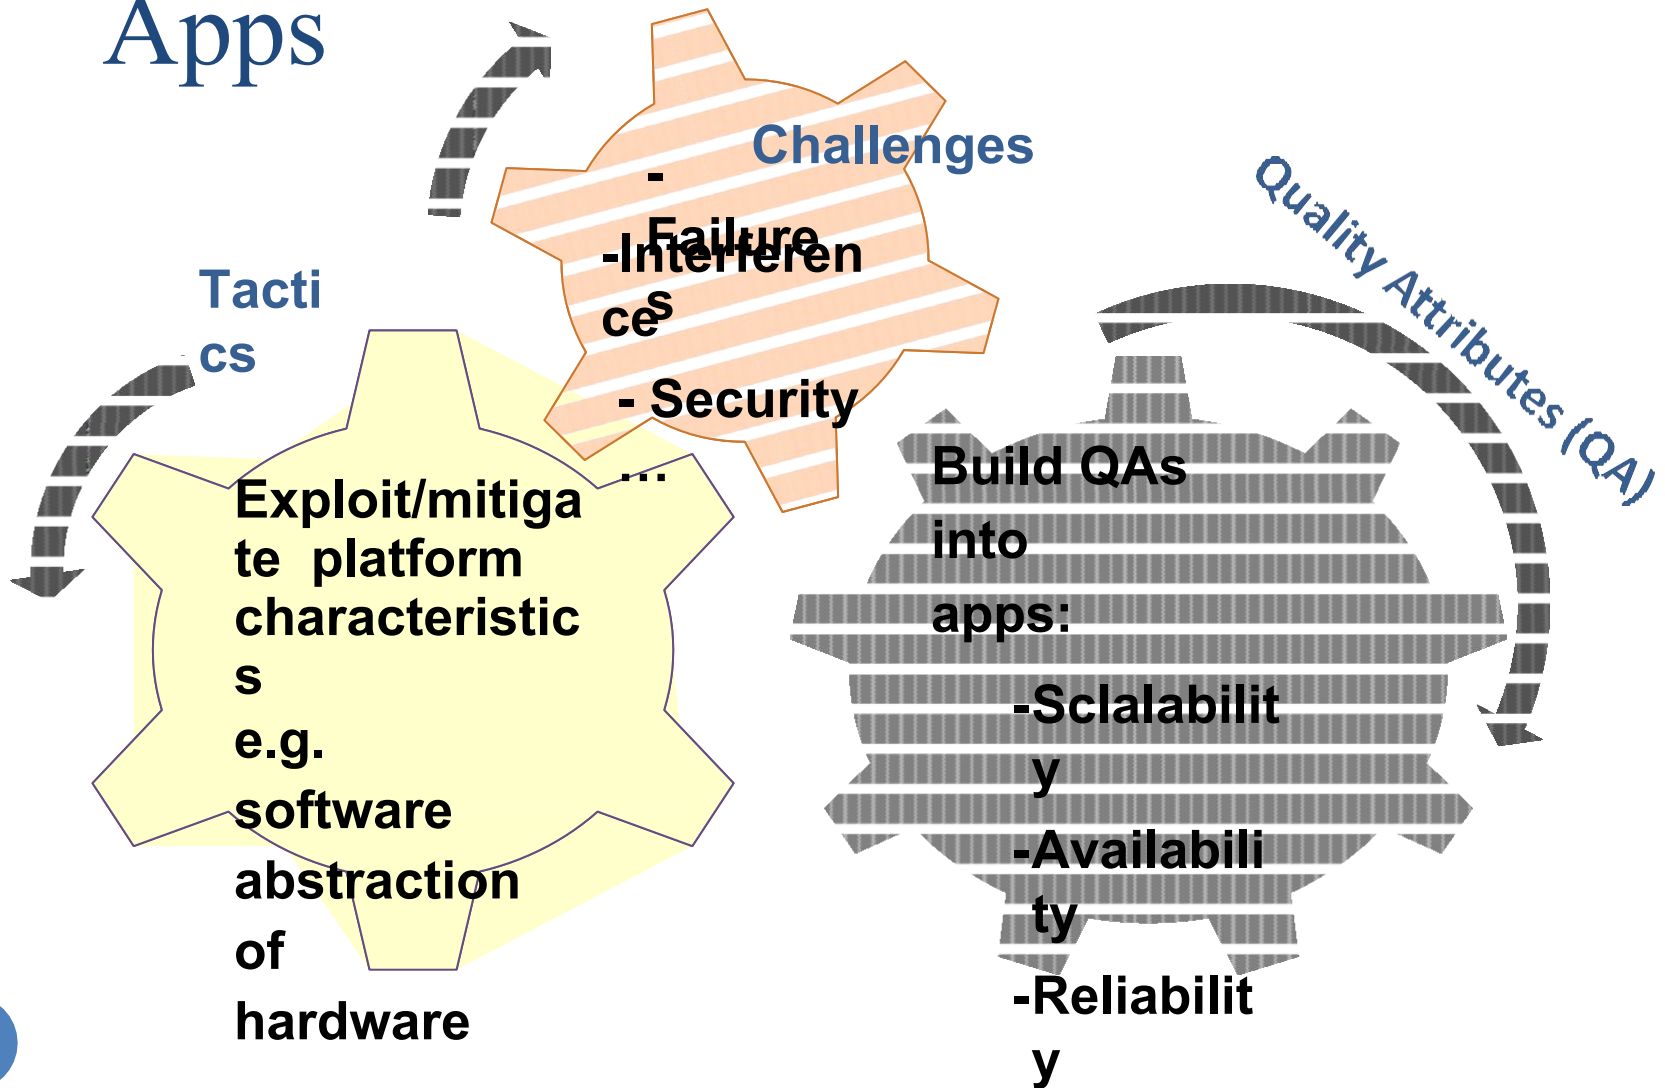


Cloud-specific Architecture Issues

Platforms & Applications Architecture



Building *Quality* Into Cloud Apps



Cloud vs. Non-cloud

- Cloud vendors address some issues faced by non-cloud data centers
 - Provisioning of server resources
 - Physical security
 - Hiring/training data center personnel
- Several other issues still remain the same
 - Network intrusion threats from outside
 - Isolating and managing production/test environments
 - Installation of updates/patches
- Some new or now-changed issues have arisen
 - These are the issues that we bring out

Key Issues Specific To Cloud

- Security/privacy
 - Multi-tenancy
 - Access keys/credentials
 - Dependency on geographic/legal jurisdiction
- Failure
 - Failure of VM instances
 - Data consistency failures
 - Software upgrade error

Key Issues Specific To Cloud

- Performance
 - Network latency
 - How fast can you provision
 - Elasticity □ Over/under provisioning
 - Performance interference due to VM co-location

Details about Key Issues

Failure Defined

- *A system failure occurs when the **delivered service no longer complies with the specifications**, the latter being an agreed description of the system's expected function and/or service*
- Related terms:
 - Faults (defects or bugs)
 - Errors (expected and actual behavior differs)
- Faults and errors **may** lead to failure

Failure Rates Get Amplified In Cloud

- Failure rates
 - Servers experience 2-4% annual failure rates (AFR)
 - Disk drives have about 4-6% AFR
- For server AFR of 3% MTBF is about 292000hrs
 - More than 30 years
- In a datacenter having 64000 servers having 2 disks each
 - Daily, more than 5 servers and 15 disks can fail

Handling Failure In Cloud

- ❑ Before you handle a failure you must detect it
- ❑ *Heartbeat* remains the key tactic for identifying failures
- ❑ Must have a *monitor* that watches aliveness of VMs
 - ❑ A monitor can be in: infrastructure, client or part of the application
- ❑ A VM must periodically show its aliveness to the monitor
 - ❑ By responding to some query
 - ❑ Sending some message

Stateful vs. Stateless Instances

- Applications can be stateful or stateless
- Stateful applications
 - Remember information across requests
 - State information kept in memory or external
- Stateless applications
 - Require that requests carry necessary information needed for understanding the context
- Failure recovery is different in stateful and stateless cases

Recovery For Stateful Instance

- Multiple tactics are possible depending upon
 - Whether application stored state in the VM itself or on external device
 - Application's tolerance for loss of computation
 - Availability of VM check-pointing
- Basic tactic is to:
 - Keep saving most recent state data somewhere safely
 - Restore back to last saved state on detecting failure


Stateful Recovery | Tactic 1

- Assumes that application keeps state in VM itself
 - No state dependency on external device
- Works for cloud only
- *Check-point VM state on regular intervals
- On detection of failure, restore and start from last check-pointed VM image
 - Requests arriving between time of failure until restoring to last checkpoint get lost

*** Sodhi & Prabhakar, “Cloud-oriented platforms: Bearing on Application Architecture and Design patterns” in IEEE Services 2012**

Stateful Recovery | Tactic

1




Stateful Recovery | Tactic 2

- Application saves recent state on external device
 - Done at regular intervals
- On detection of failure:
 - Check for recent state saved on external device
 - Resume from state restored from external device
- Requests arriving between time of failure until restoring to last saved state get lost
- Recovery mechanism needs to be coded into application itself

Stateful Recovery | Tactic

2



Avoiding Lost Requests In Tactic

2

- Application logs incoming request as well to external device
- On detection of failure:
 - Resume from state restored from external device
 - Read and play logged requests that arrived during down time
- No requests are lost if logging to external device and acknowledgement to client is atomic

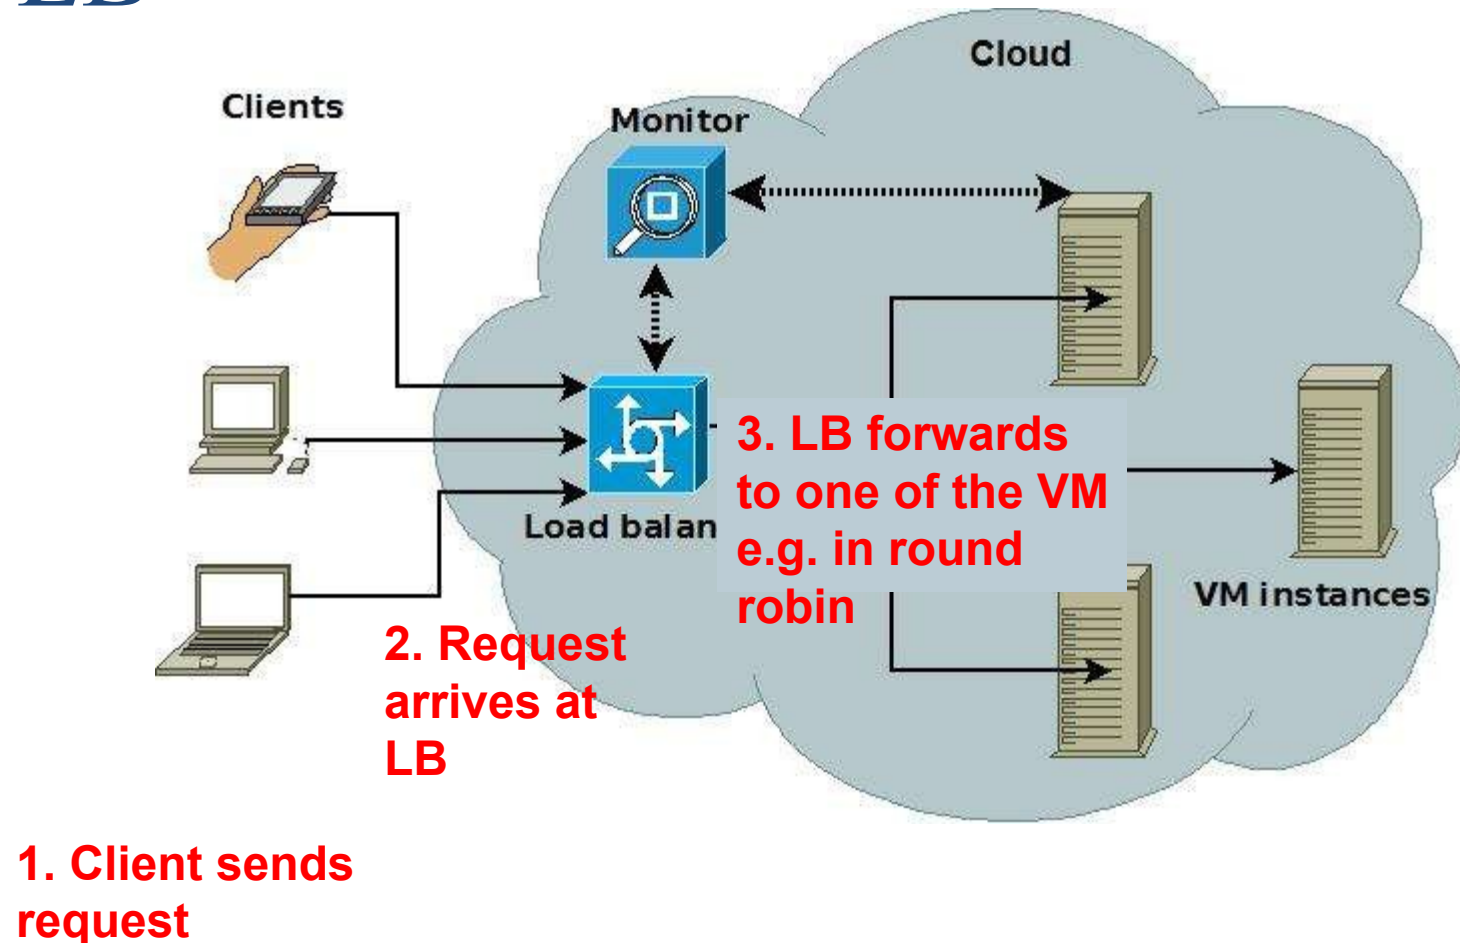
Stateless Applications | Requests Flow

- A request can be received by any instance
- New VMs can be created to meet QoS demands
 - Performance, reliability etc.
- Requests are routed to instances by a load balancer component

Routing Requests To Instances

- Decided by load balancer (LB)
 - LB policies take failures into account
 - LB is augmented by a monitor component
- Tactic has two flavors:
 - **Push**: LB decides which instance gets to serve a request
 - **Pull**: Instances pull the requests from a queue maintained by the LB

Architecture Of Push Based LB



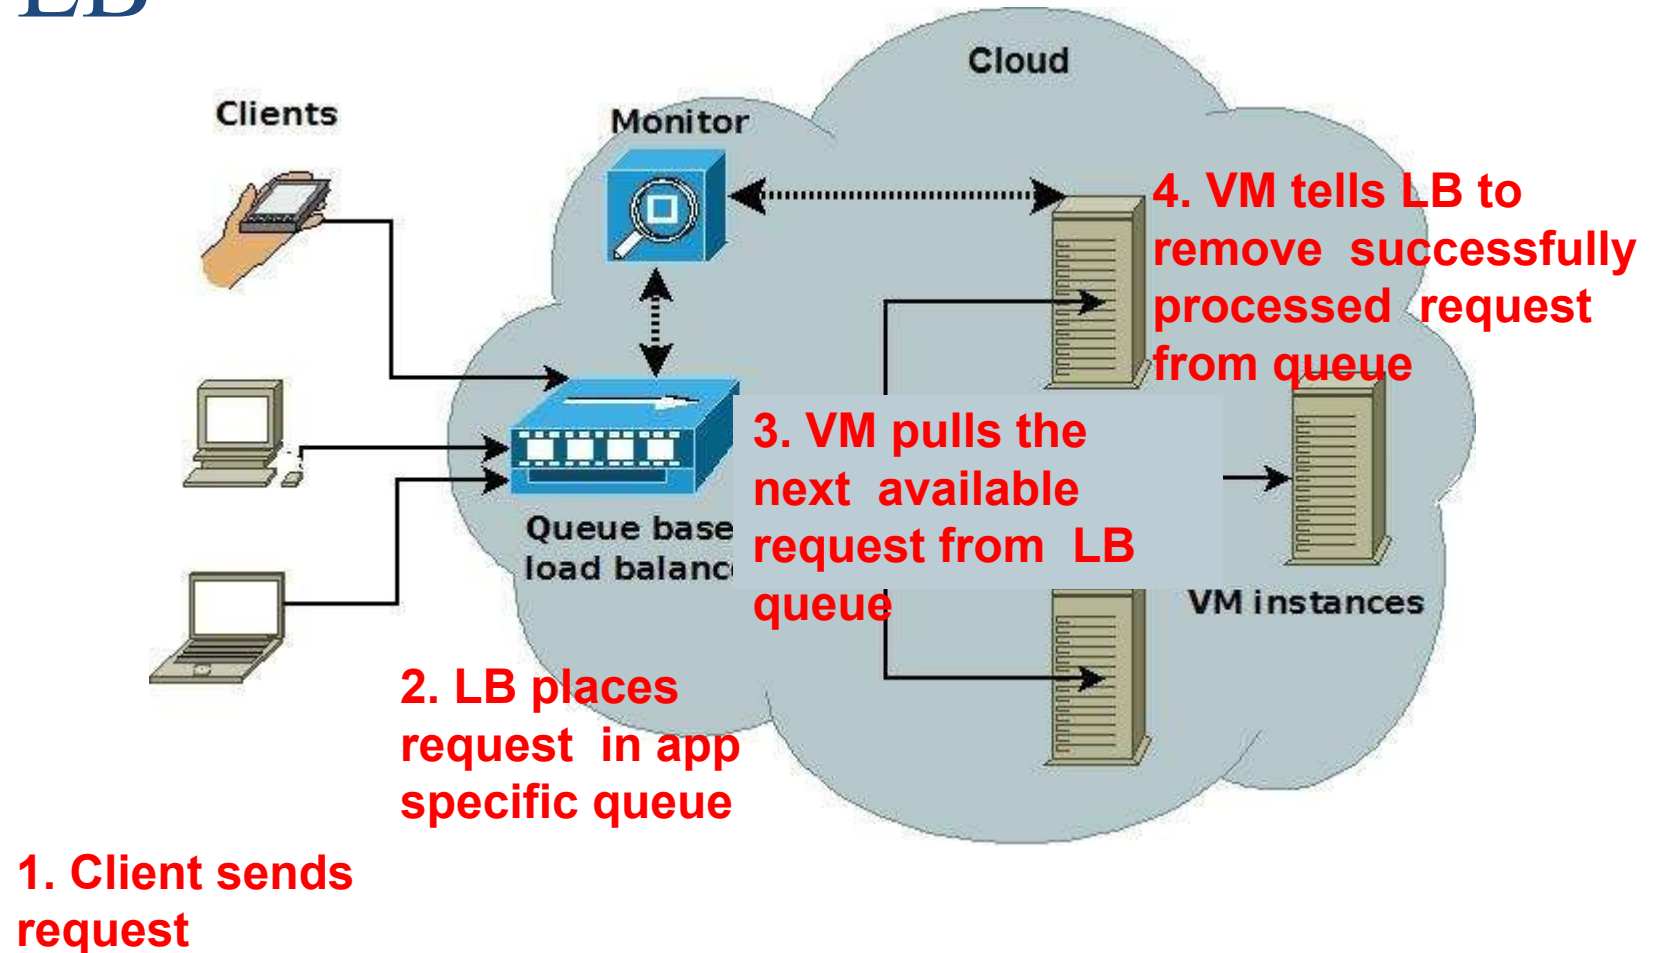
Role Of Monitor

- Observes the VMs
 - Resource (CPU, memory etc.) utilization
 - Requests load
 - Quality of service violations etc.
- Sends VM stats to load balancer
 - Include failures info
- Decides, based on some rules, when more resources are required

Working Of Push LB Pattern

- Monitor detects VM failure
 - E.g. when VM becomes non-responding
- LB gets to know this and stops sending requests to the failed VM
- Current in-progress requests handled by VM are lost
 - Client needs to detect this possibly via timeout
 - Client should resend the lost requests

Architecture Of Pull Based LB



Role Of Monitor

- Watch the application specific queues on LB
 - Waiting time for requests in a queue
 - Length of the queue
- Infer load on a VM from queue stats
- Decides, based on some rules, when more resources are required

Working Of Pull LB Pattern

- LB knows when a request has been processed by a VM
- Requests that remain unhandled until a time limit get reassigned by LB
- A failed VM won't pick requests from its queue
 - This automatically takes a failed VM out of service
- Requests trapped in failed VMs
 - Can get processed when VM recovers
 - Application must handle duplicate processing scenario

VM Cleanup On IaaS Cloud

- When a VM fails:
 - It is not automatically de-allocated
 - It needs to be de-allocated by consumer
 - Cloud provider continue to charge until VM is de- allocated
- On de-allocation of a VM:
 - Its public and private IP addresses become available for reassignment
 - Infrastructure can be told to assign released public IP address to replacement VM

Summary

- LB and monitor are key components
 - LB policies take failures into account
 - LB is augmented by a monitor component
- Tactic has two flavors:
 - **Push**: LB decides which instance gets to serve a request
 - **Pull**: Instances pull the requests from a queue maintained by the LB

Data storage issues on cloud

Data Storage

- Provides ability to persists digitized information
 - e.g. plain text files, binary files containing photos
- Retains data for an interval of time
 - Length of time depends on storage type
 - e.g. RAM contents don't survive machine restart, whereas hard disk contents can

Data Storage Components

- Can be have different categories, for example:
 - Raw files on a file systems (FS) on the operating system
 - Data stores such as RDBMS engines, key-value stores etc.
- Characteristics/behavior of each category can vary significantly
- An application can use multiple types of data storage components
 - E.g. can write to both the raw FS and a database table

Data Storage And Cloud

- Mainly we are concerned about data storage on:
 - IaaS cloud
 - PaaS cloud
- IaaS cloud because:
 - Cloud user has to manage the resources including storage
- PaaS cloud because:
 - As a developer you need to handle data storage from within the applications you write
- SaaS cloud case is not interesting
 - Because as a user you don't write any software here

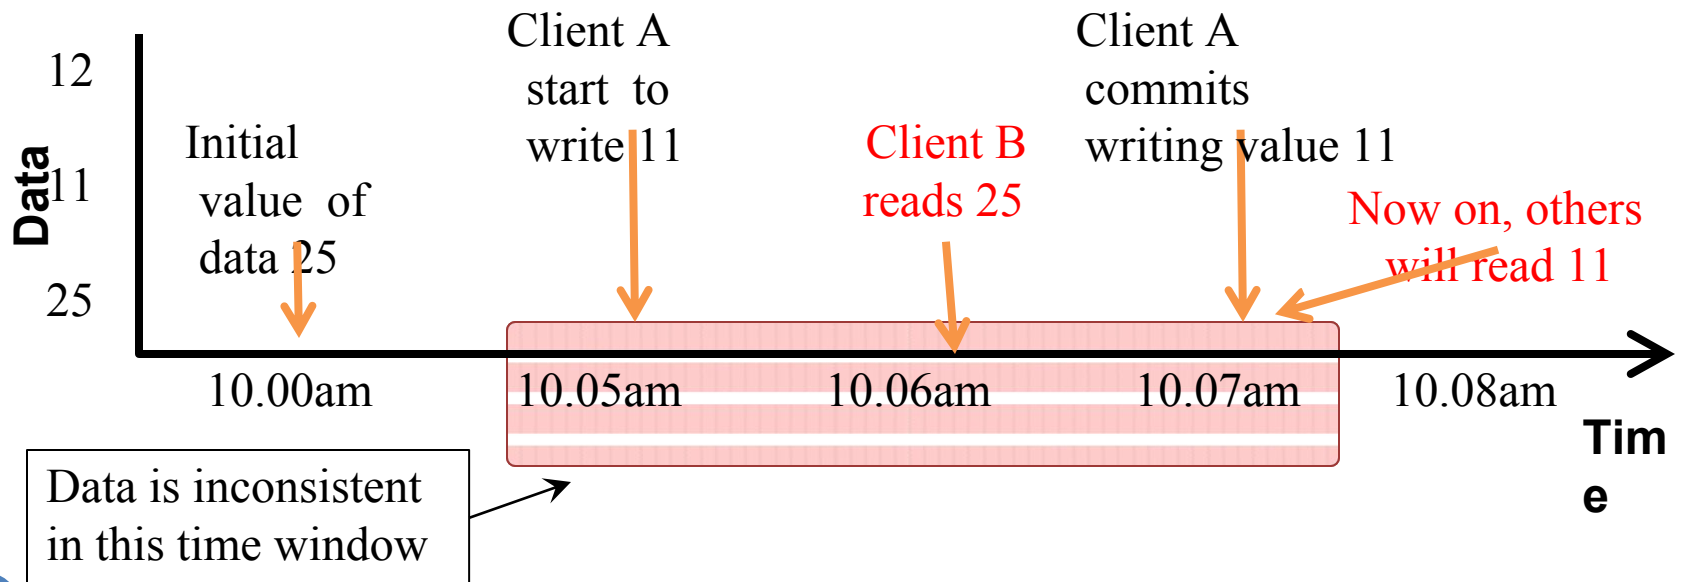
Data Storage On Cloud

- ❑ IaaS cloud vendors offer two data storage types:
 - ❑ Ephemeral
 - ❑ Persistent
- ❑ Ephemeral storage doesn't survive instances failure
 - ❑ Typically available as a block device attached to the VM instance
- ❑ Persistent storage is long lived
 - ❑ Cloud vendor automatically replicates
 - ❑ Geographically distributed replicas
- ❑ Storage failures can lead to data consistency issues
 - ❑ Application needs to be prepared for this

Data Consistency In Applications

□ Consistency

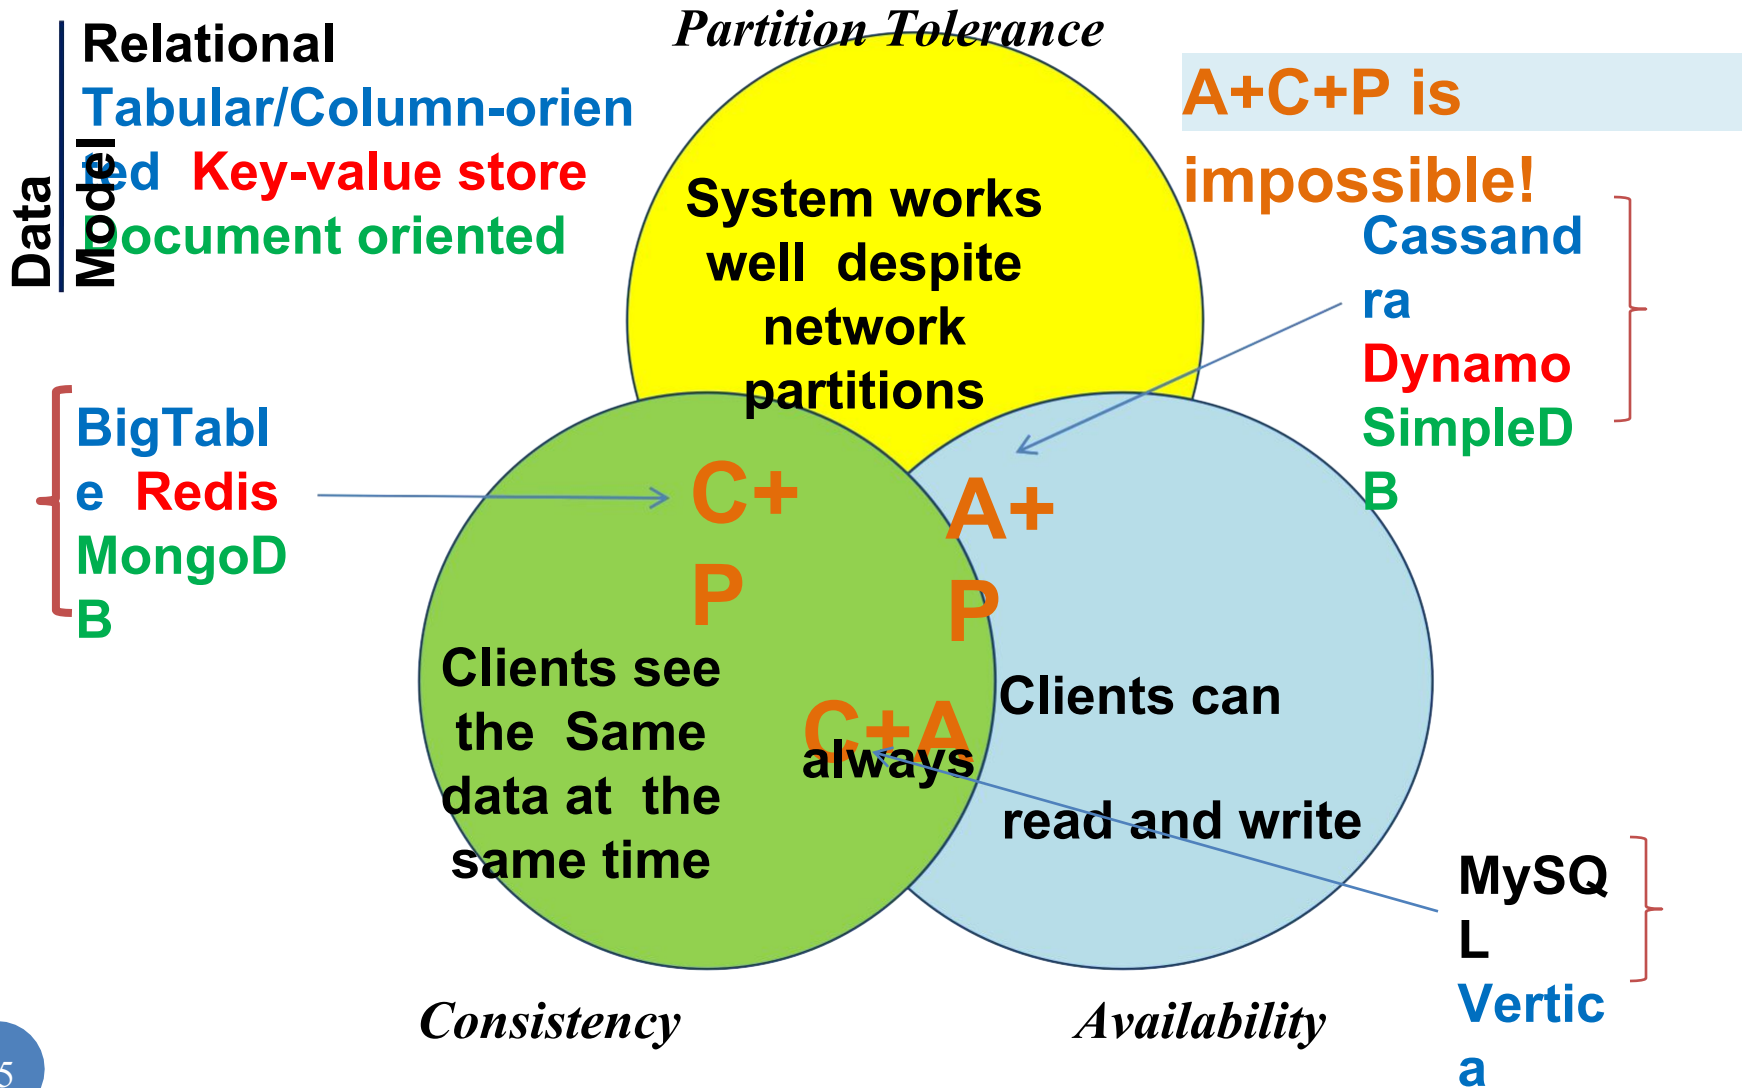
- Disallow multiple values of same piece of data when seen by different clients at the same point in time



Data Partition Tolerance

- Partitioning of data means:
 - Placing parts or copies of data on multiple nodes
 - A “node” here may mean a physical machine or just a different database server instance
 - Often needed to achieve massive scalability
- Tolerance for data partitioning means:
 - Application should not respond incorrectly even in the presence of data partitioning
 - i.e. data *consistency* is maintained in presence of partitions

Data Consistency | CAP Theorem



Data Consistency In Cloud

- Replication is a key mechanism
- Replicating data is time consuming
 - Replicas may not be in sync immediately
 - Clients may see inconsistent data for a tiny duration
- Typically, cloud vendors offer “consistent reads”
 - But there may be latency issues
- E.g. AWS S3 today offers read-after-write consistency for PUTs of new objects
 - But gives only eventual consistency for overwrite PUTS and DELETES

Why Is It Important?

- Because users needs to be happy
 - “500 Internal Server Error” is the last thing I want to see after punching in my credit card details
 - Next time I’ll shop elsewhere
- It impacts* the businesses too
 - Extra 0.1s in response time costs Amazon 1% in sales
 - Google found that 0.5s jump in latency leads to 20% drop in traffic

* <http://highscalability.com/latency-everywhere-and-it-costs-you-sales-how-crush-it>

How To Address It

- You need to choose any two from among:
 - Availability, consistency and partition tolerance
- Drop availability
 - Services are unavailable until data is consistent on all nodes
- Drop partition tolerance
 - Avoid partitions from happening
- Drop consistency
 - Expect that data becomes consistent eventually

Summary

- It is important to understand characteristics of data storage mechanisms on cloud
 - Different storage services may behave differently
- Data consistency in applications
 - Cannot achieve all three at the same time:
 - Consistency of data
 - Availability of data/services
 - Partition tolerance

Software Upgrade Induced Failures

- Upgrading and patching software is common
- Main reasons for upgrades and patching
 - Releasing new features
 - Fixing defects
- Users may experience failures during such upgrades and patching of software applications
 - How to avoid such failures?

An Example Scenario

**Client/web
browser**

**Server-side
VM**

**t2) Client initiates
request**

**t1) S/W upgrade
starts**

**t4) Server
reply**

**t3) Upgraded version
of application
handles the
request**

**t5) Subsequent
request in the
session**

**t6) Old version of
application
handles the
request**



Handling S/W Upgrade Failures

- What do we need from the solution
 - Clients should always interact with latest application version on server side
 - Requests should get load balanced regardless of active application versions on server side
- We need to make some assumptions:
 - Backward compatibility of application versions
 - Clients can be dynamically told about latest application version number available

A Solution Idea

- Treat each application version as a separate destination for requests
- Client knows the latest version number that it has interacted with
- Load balancer routes requests based on version number present in request header
 - A version xxx in the header is routed to instance whose version is \geq xxx
 - Route normally if no version number found in request header

Security and privacy

Information Security

- “Information security means protecting information and information systems from unauthorized access, use, disclosure, disruption, modification, perusal, inspection, recording or destruction”

(From 44 US Code § 3542)

<http://www.law.cornell.edu/uscode/text/44/3542>

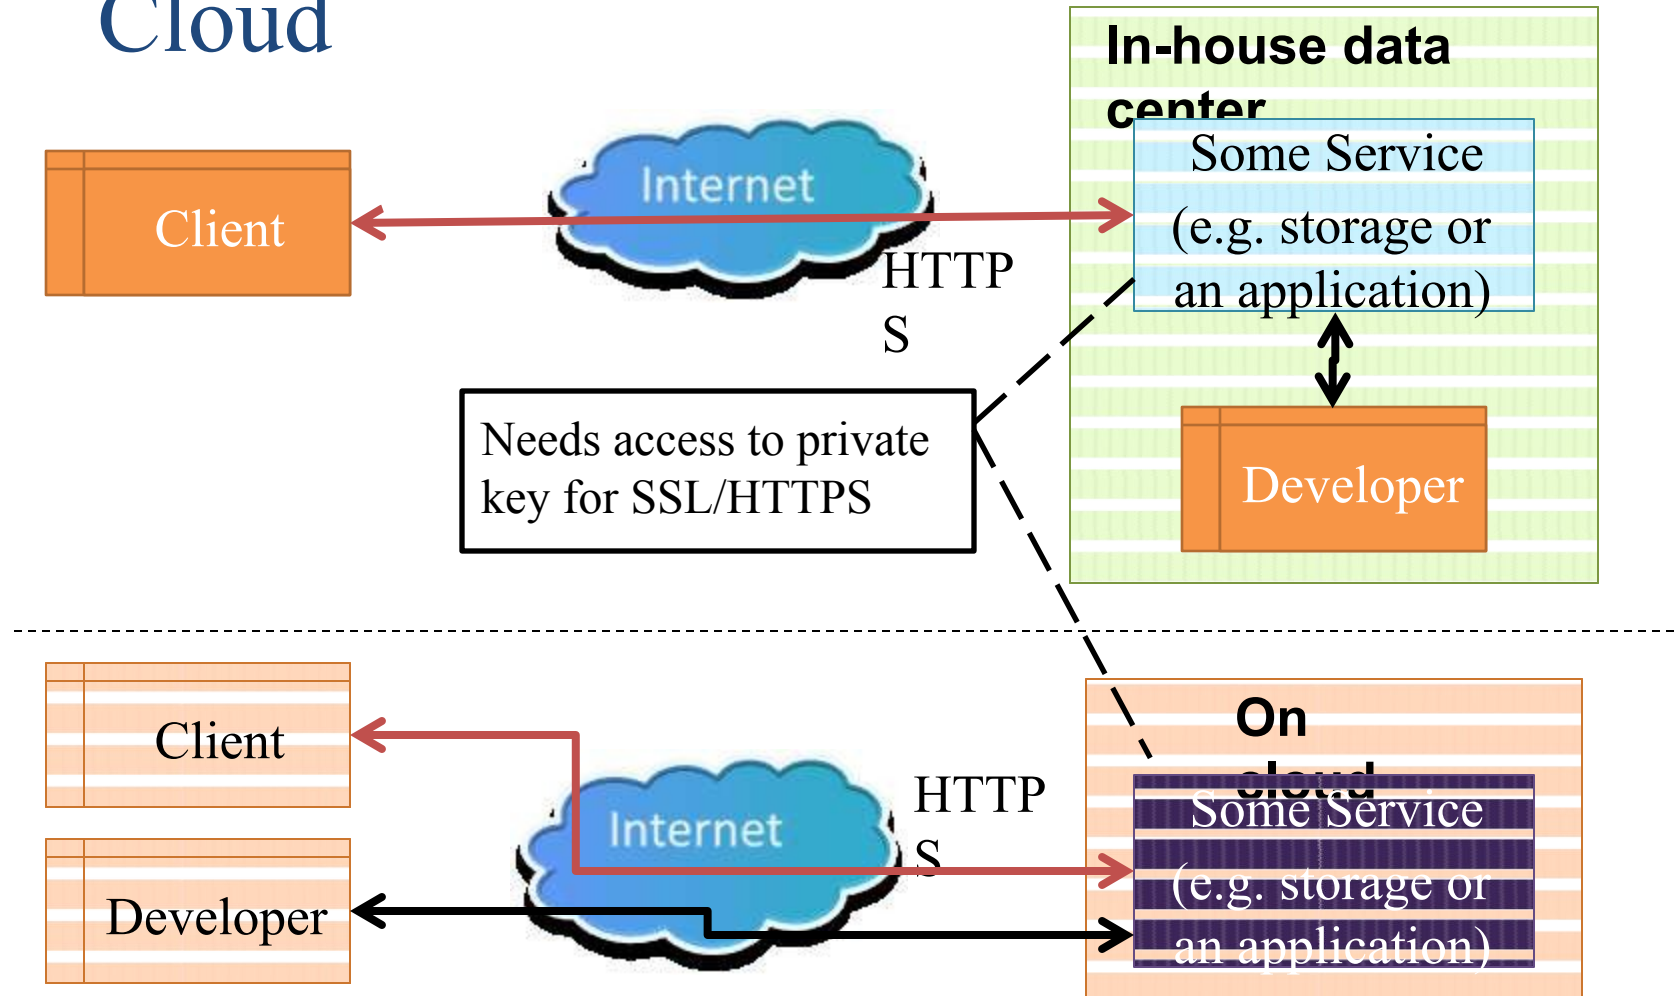
Core Artifacts: Credentials and Keys

- **Credentials** in IT are typically, a combination of login ID and password
 - Identify the user holding them
- Mainly credentials are used for controlling access to IT resources
 - **Authentication**: you are who you say you are
 - **Authorization**: you have the rights to perform certain actions
 - **Non-repudiation**: you cannot deny you did something
- A **key** is a number used in cryptography for
 - Encrypting/decrypting data
 - Digital credentials

Some Common Security Measures

- ☐ Encrypting sensitive data
- ☐ Checking for buffer overruns
- ☐ Input validation
- ☐ Role based access to data and functionality
- ☐ Transaction monitoring and auditing
- ☐ Limiting access to servers (close unnecessary ports)
- ☐ etc.

Security Scenario | In-house vs. Cloud



Important Security Aspects On Cloud

- Central issue:
 - Data and applications being in 3rd party (cloud vendor) custody
 - Lack of trust on the cloud service provider
- Access credentials and encryption keys
 - Management of keys and credentials
- Privacy and security in multi-tenant environment
- Dependency on legal/geographical jurisdiction
 - Local laws governing a cloud service provider

Bank vs. At-home | A Metaphor


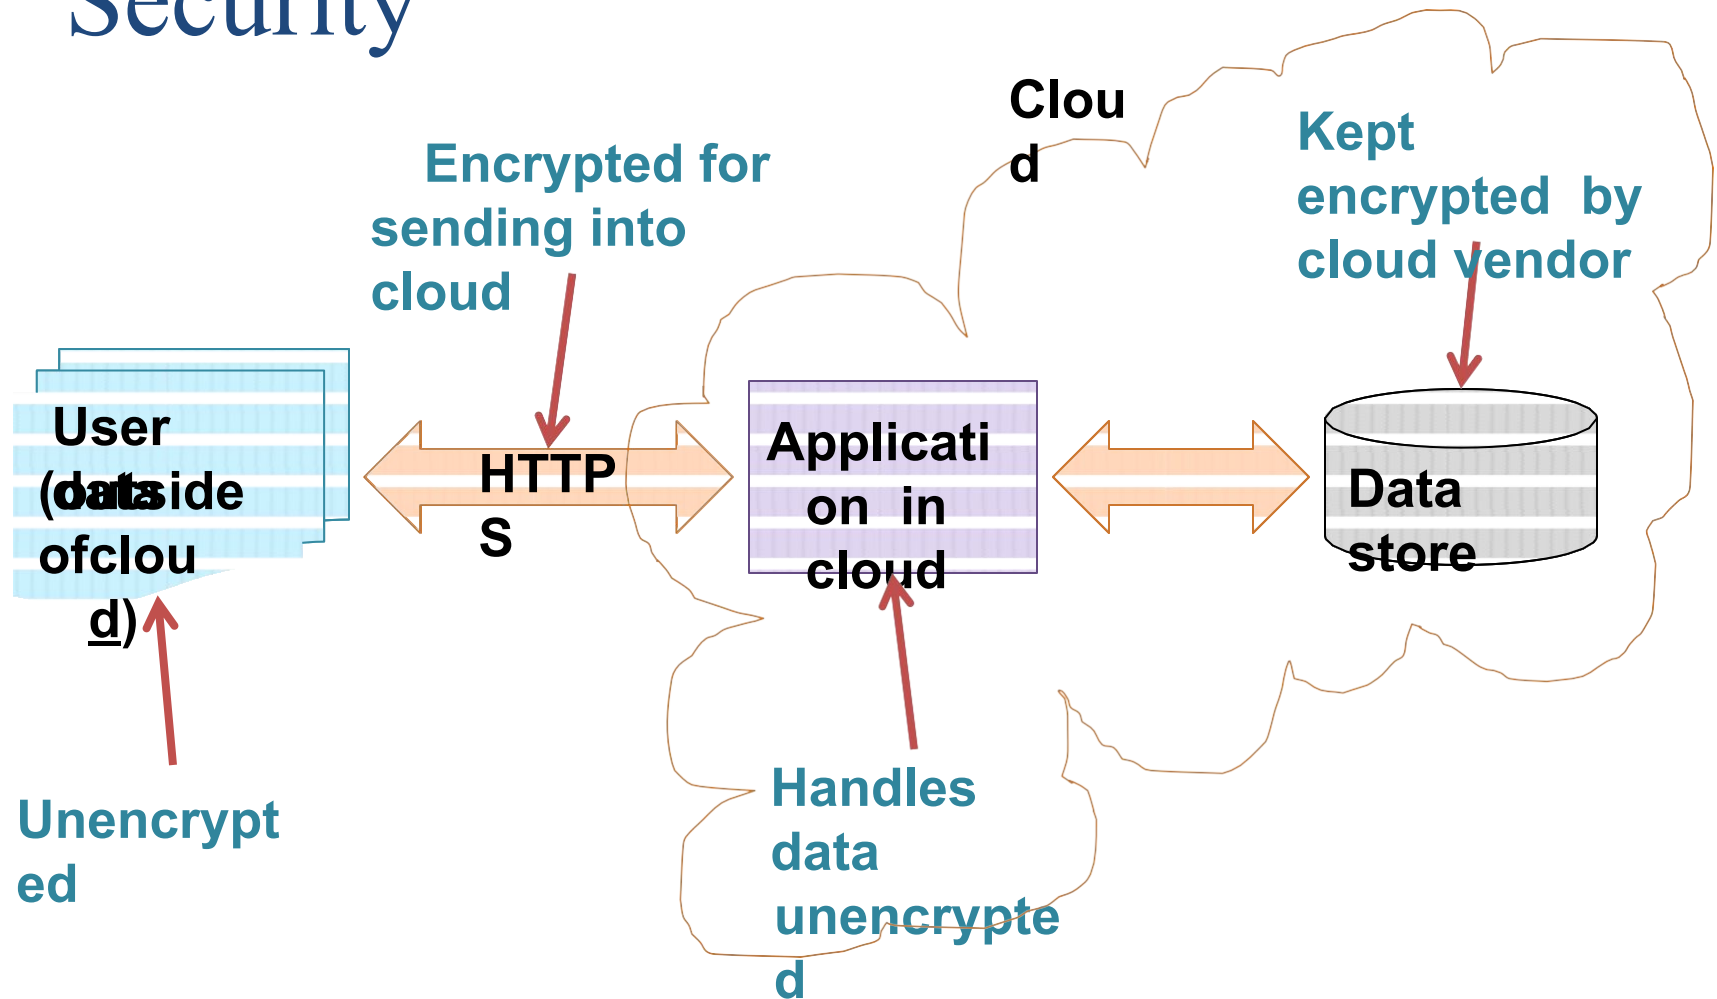


Image credit: Business Today

Elementary Data Security



Issues With Elementary Data Security

- A cloud provider responding to a legal order may:
 - Potentially share your data with govt. agencies
 - Have to provide decrypted data
- To address the above you can encrypt your data before sending to cloud
 - Thus cloud provider can share only encrypted data
- However, if legal orders are directed at you then you need to comply with decrypted data

Credentials Management For Cloud

- Important considerations:
 - Who has/needs access to credentials?
 - Will you need to change credentials? When?
 - Storage and automated provisioning of credentials
- Some options for providing credentials in cloud
 - Build into the VM image
 - Supply as parameter during instance launch
 - Keep in some persistent storage
 - Send from client every time a new instance starts

Summary

- Several security aspects remain same on cloud as they were in-house
 - e.g. from tightening of firewall rules to input validation in you application
- Few things have changed
 - Data/apps now lives in a 3rd party custody
 - It is cloud consumer's responsibility to protect its sensitive data
 - Can use encryption technologies to achieve this

Geo/Legal Jurisdiction Dependency

- Requirements related to storage of personal information (PI)
 - The EU mandates that PI can leave the EU only for jurisdictions that guarantee equivalent privacy
- Some jurisdictions claim access rights on all data stored inside their borders
 - US Patriot Act allows any data stored in US to be examined by US govt.

Impact On Cloud Consumers

- Awareness of cloud vendor data centers location
 - Some do not provide locations of its data centers
- Backup locations may be chosen by the vendor
 - For optimization of its resources etc.
- Abilities to allow users to control data locations
 - Based on type of the data
 - Every vendor may not offer

How To Deal With It?

- Use anonymization techniques for securing PI
 - E.g. replace PI with tokens and keep the **PI ↔ token** mapping locally
- Example:
 - Original data:
 - Shiva Kumar □ {Sensitive data}, e.g., bank account number
 - Anonymization tokens (kept locally)
 - Shiva Kumar □ {Token}, e.g., a number
 - Data stored in cloud:
 - Token
 - Sensitive data
 - Restore original data by taking join of token table and cloud data table

performance

Performance Of A System

- Characterized by quantum of **useful work** done by a system **for a given amount of time and computing resources**
- Often measured by
 - Response time (**smaller the better**)
 - Latency (**smaller the better**)
 - Throughput (**higher the better**)
 - Resource utilization (**higher the better**)

Achieving Better Performance

- System should do:
 - more work
 - at a faster rate
 - by consuming less computing resources
- Time-proven design principles apply on cloud as well
 - Exploit parallelism
 - Pooling of shared resources
 - Put processing near the data/resources it needs
 - Minimizing round trips
 - ... and rest of the good stuff

Key Points For Cloud

- Consolidation of computing resources
 - Improves overall utilization in a data center
 - Use virtualization to achieve this
- Rapid elasticity
 - On-demand provisioning of resources
 - Fast scaling of applications
 - Latencies can still be an issue
- Multi-tenancy
 - May result in performance interference

Elasticity On Cloud Platforms

- Elasticity: provision computing resources on demand
- Consumers can define auto-scaling rules
 - E.g. on an existing VM when CPU usage remains above 80% for 10 minutes, launch a new VM
 - Eliminates manual intervention for provisioning
- Auto-scaling strategies
 - A matter of research
 - Provisioning a VM and starting the apps on it takes time

Provisioning Latency: An Important Issue

- Small Instance

- 1 EC2 Compute Unit (1 virtual core with 1 EC2 Compute Unit),

- 1.7 GB of memory, 160 GB of instance storage,
32-bit platform with a base install of CentOS 5.3
AMI

- Can take **5 to 6 minutes** us-east-1c from launch to availability

- Large Instance

- 4 EC2 Compute Units (2 virtual cores with 2 EC2 Compute Units each), 7.5 GB of memory, 850 GB of instance storage, 64-bit platform with a base install of CentOS 5.3
AMI

- Can take **11 to 18 minutes** us-east-1c

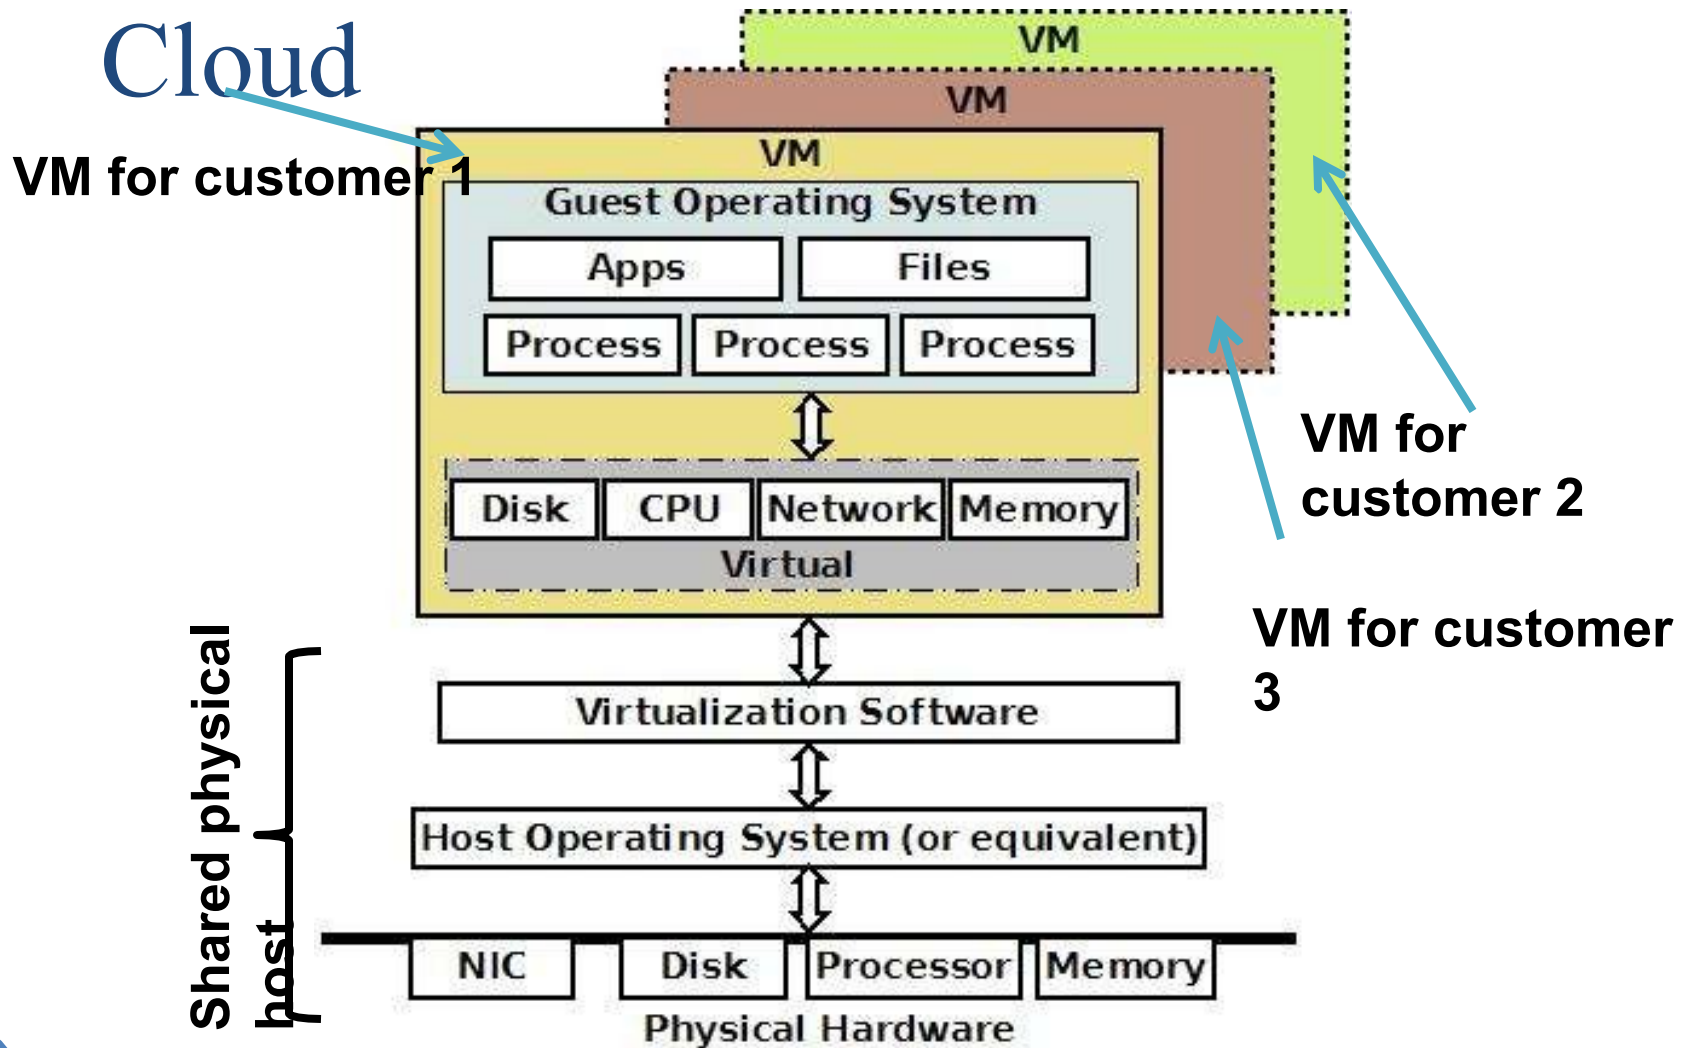
Addressing Provisioning Latency Issue

- You need to forecast the number of resources needed for optimal service
- Basic idea is to:
 - Watch for events that may indicate potential surge in load
 - Calculate the amount of additional resources needed to handle load
- Success depends on:
 - How well you pick right events that indicate load surge
 - Accuracy of calculating needed resources

Multi-tenancy On Cloud

- Different consumer's applications and data hosted on shared infrastructure
 - e.g. Single physical disk holding data from different consumers
- Needed for optimizing resource utilization
 - Allows cloud provider to leverage economies of scale
- Consumers can assume they are sandboxed
 - i.e. their apps and data is isolated from other's

Multi-tenancy On Cloud




Issues Arising Due To

Multi-tenancy

- Performance interference
 - VMs co-located on same physical host may affect each other's performance
 - E.g. simply forcing expensive cache invalidation on the host CPU
- Potential for one VM breaking into another
 - Bugs in virtualization software
 - Encrypting data by user can help

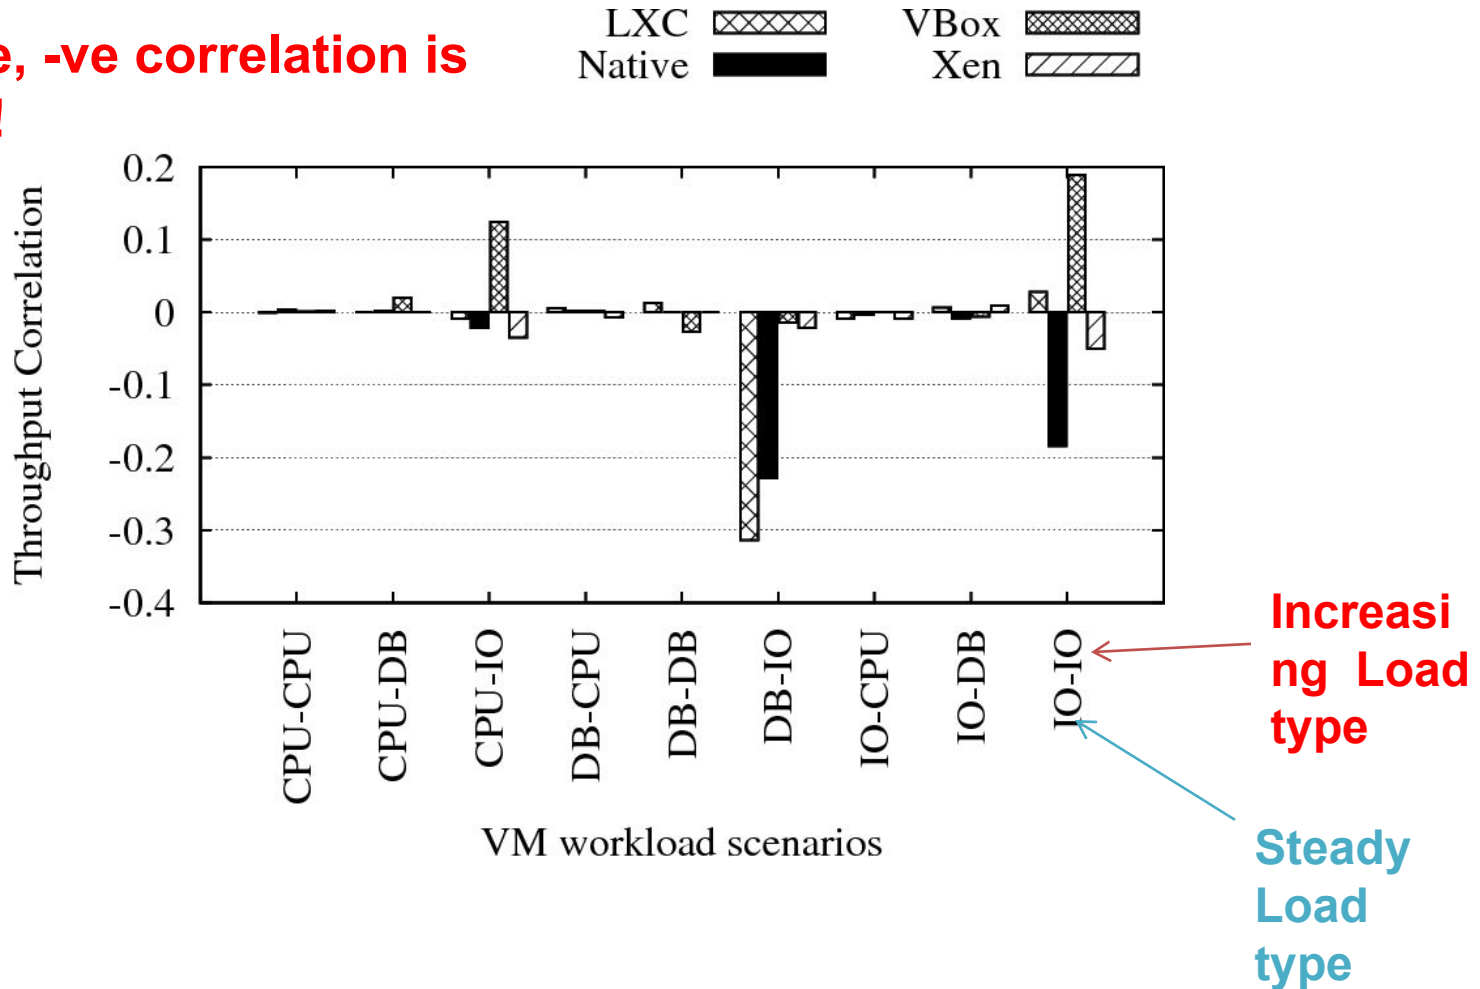
Sodhi & Prabhakar, "Cloud Platforms: Impact on Guest Application Quality Attributes" in IEEE APSCC 2012

Measuring Interference




Performance Interference (Throughput)

Here, -ve correlation is bad!




Performance Interference (CPU)

Here, +ve correlation is bad!



Performance Interference (Memory)

Here, +ve correlation is bad!



Summary

- Performance best practice from non-cloud environments continue to apply on cloud
- But you can leverage cloud specific characteristics
 - Elasticity of computing resources
 - Virtual nature of resources
 - On-demand provisioning
- Performance remains a design issue on cloud as well
 - One has to design performance into the application
 - It does not come automatically!

Impact Of Platform Characteristics

Platform Characteristics

- A computing platform (aaS) carries characteristics unique to it
 - Chars. = {Functional + non-functional attributes of platform}
- For example:
 - Software abstraction of hardware resources (V12N)
 - Coarse-grained multi-tenancy (IaaS/PaaS)
 - Limited control of underlying infrastructure (PaaS)
 - Location transparency (XaaS)
- They impact guest application architecture
 - E.g ability of a guest app to achieve certain QAs

Finding Impact on Application Architecture

- Examine platform characteristics in light of architecture knowledge and reverse-engineer the tactics and patterns

| QA (Possible tactic for it) | Remarks |
|-----------------------------------|--|
| Reliability (State checkpoint) | On finding a failure, the system can be restored to a prior check-pointed consistent state. |
| Performance (Add concurrency) | Processing different sets of tasks in parallel by creating additional threads. |
| Security (Limit exposure) | Services are deployed on hosts in a manner that reduces overall damage when the host is compromised. |

Sodhi & Prabhakar, "Cloud Platforms: Impact on Guest Application Quality Attributes"

in IEEE APSCC 2012

Characteristics' Impact on

QAs
Some examples:

| Characteristic | Impacted QA |
|--|-----------------------------------|
| Software abstraction of hardware | DR, Efficiency (+) |
| Programmatic self-serviced provisioning | Deployability, Scalability (+) |
| VM/resource check-pointing and snap shots | Availability, DR, Reliability (+) |
| Lack of computing assets custody | Privacy, Security (-) |
| Relative anonymity behind subscription and usage | Security (-) |

Sodhi & Prabhakar, “Cloud Platforms: Impact on Guest Application Quality Attributes”

in IEEE APSCC 2012

How To Make Use Of This?

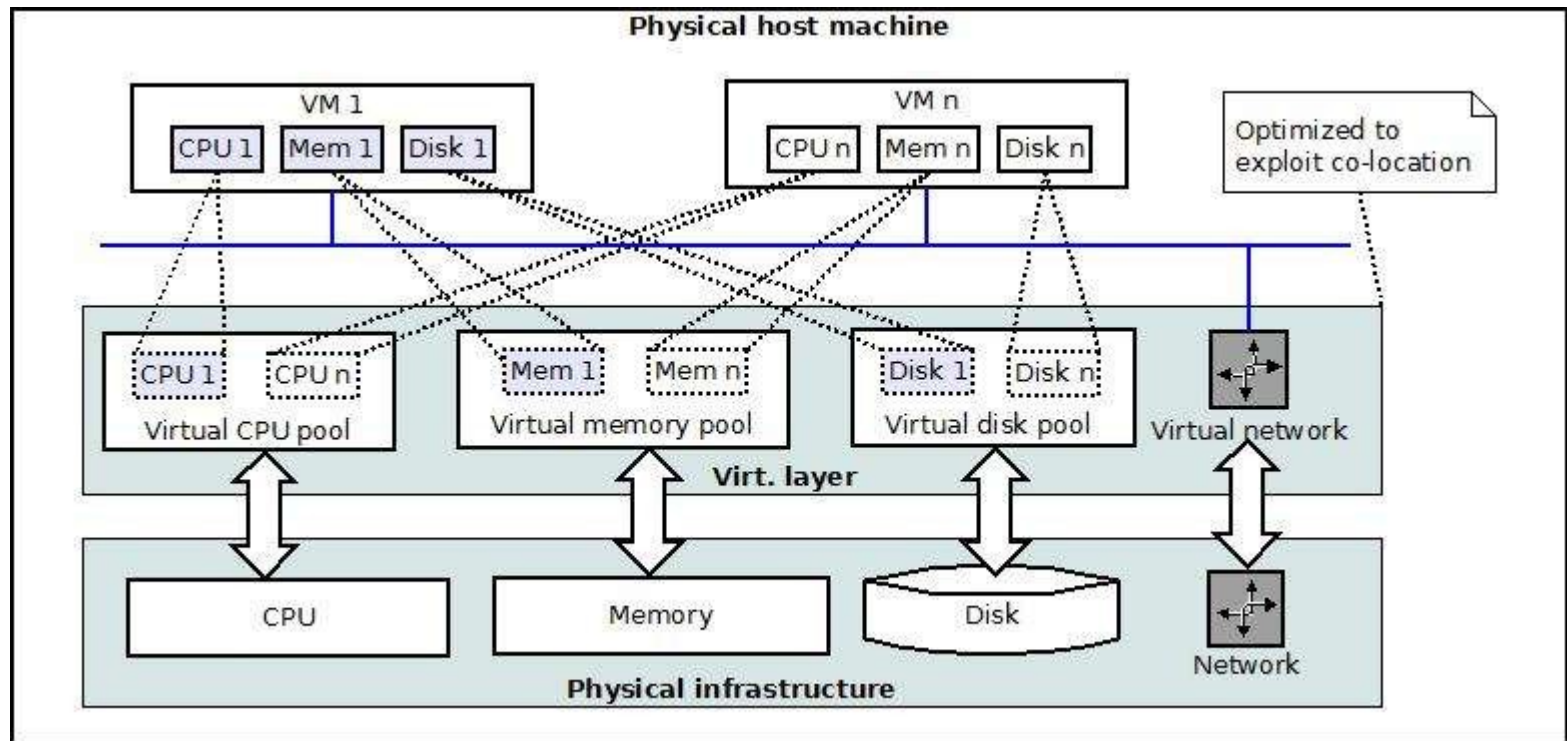
- Can be used address different types of application scenarios
 - Implement failure recovery at platform level
 - Co-locate different tiers on same host keeping them isolated
 - Selective request/computation transfer

Sodhi & Prabhakar, “Cloud-oriented platforms: Bearing on Application Architecture and Design patterns” in IEEE Services 2012

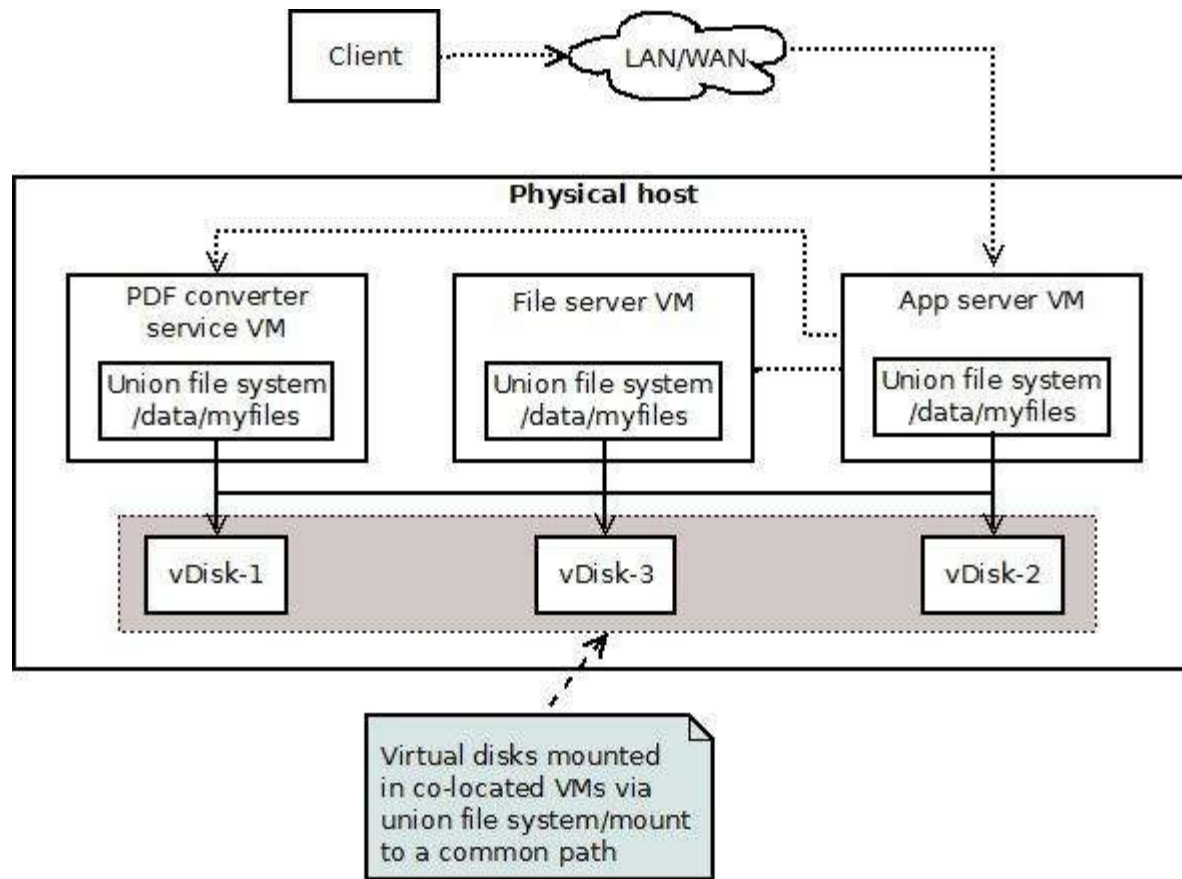
Tier Co-location: Problem Scenario

- Application consisting of multiple tiers
 - e.g. a N-tier web application
- There is some over-the-wire data exchanged among tiers
 - This is costly!
- You want to retain tier isolation, but still minimize inter-tier communication overheads

Tier Co-location: Logical View



Tier Co-location: Example

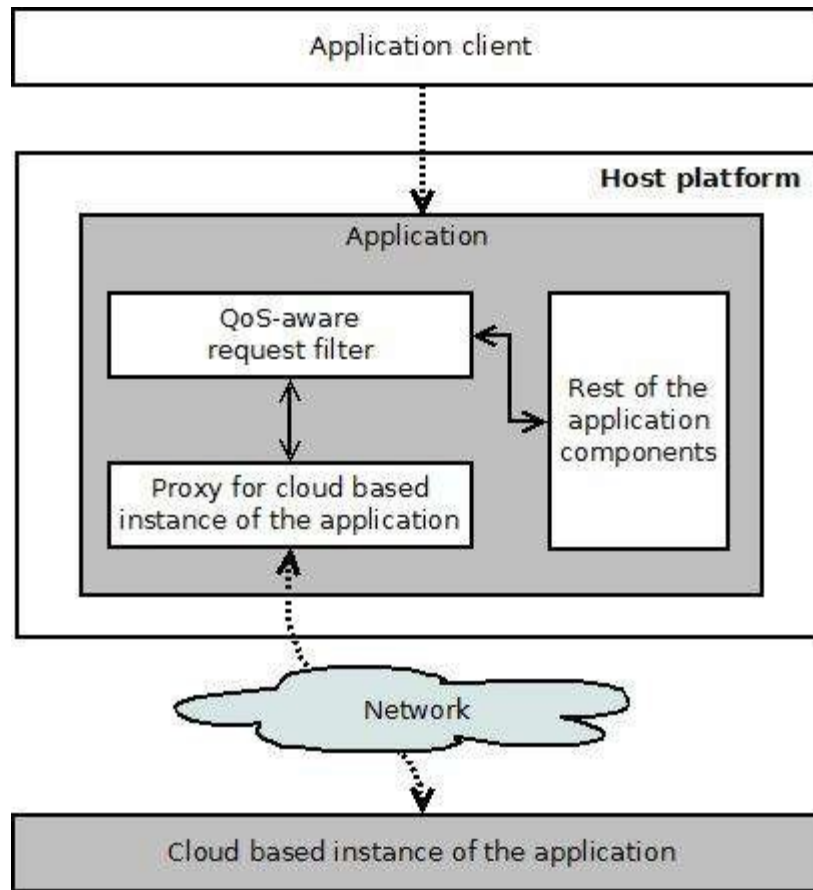


Selective Computation Transfer

- Problem context:
 - An application functionality is such that:
 - Different requests may take different amount of resources to process
 - Some requests may require execution of privileged components
 - Majority of requests can be served from normal deployment environment
 - You want to maintain certain QoS for *ALL* requests

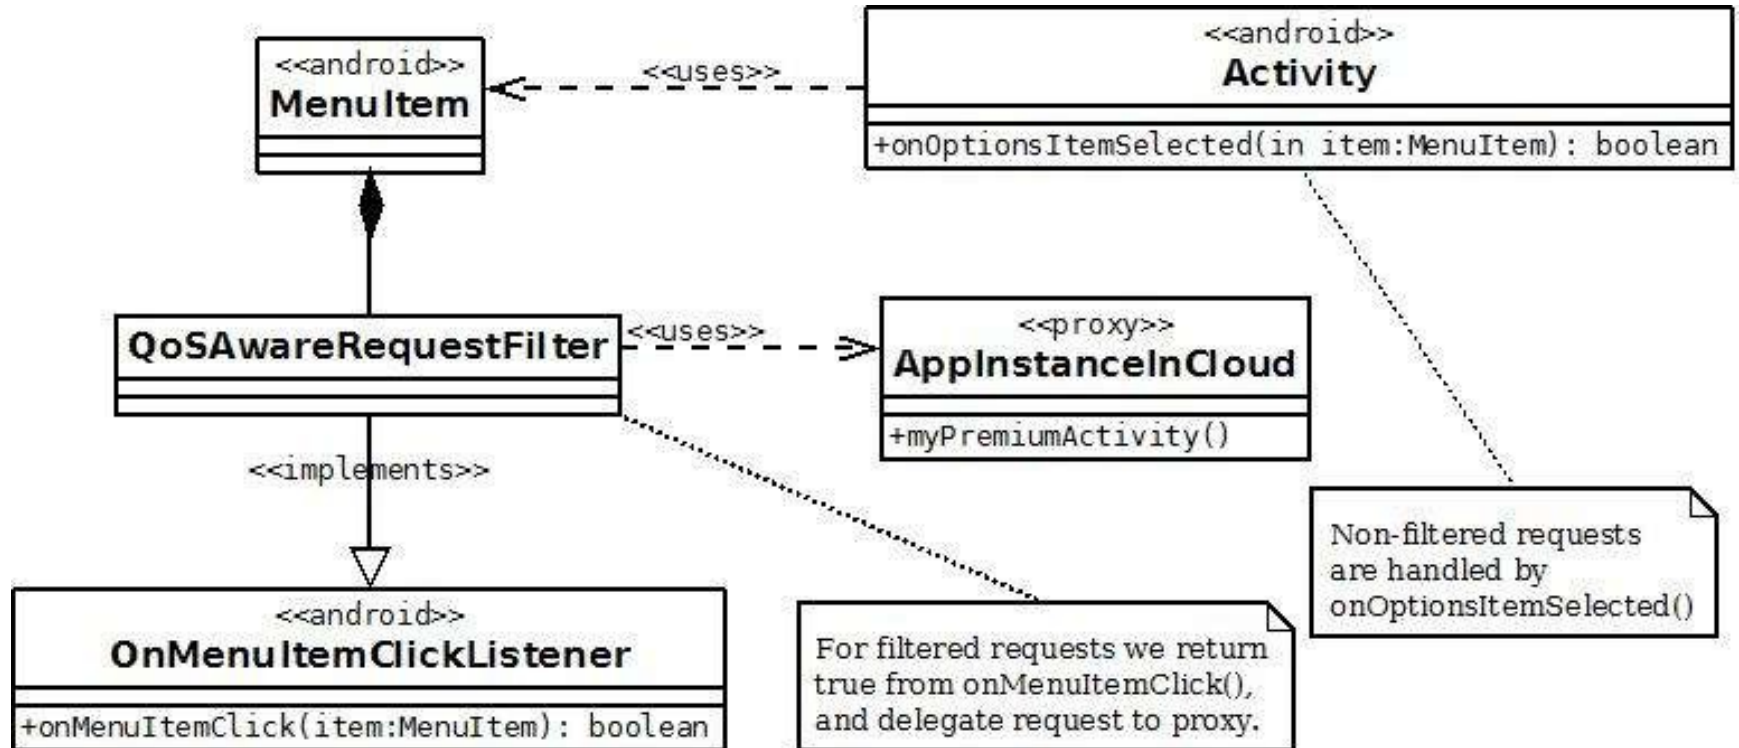
Selective Computation Transfer

Architecture



Selective Computation Transfer

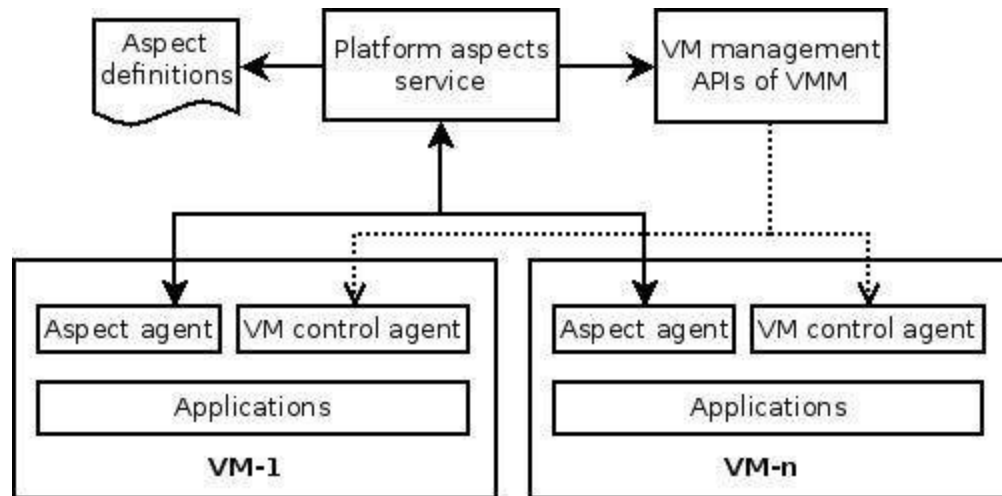
- Example implementation



Aspect-orientation at Platform-level

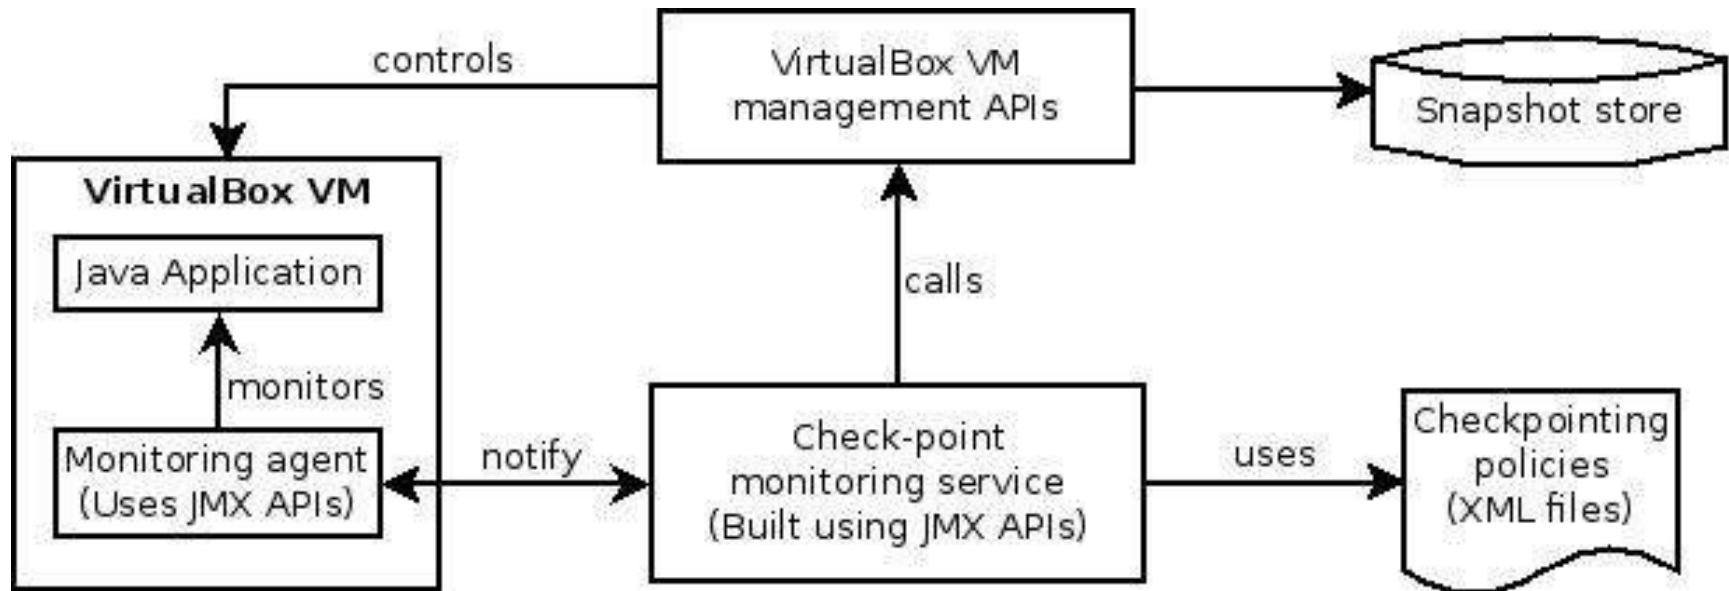
- Problem context:
 - There are software design and implementation concerns which apply at the coarse grained computing environment/platform level, and they cut across the applications
 - E.g. monitoring **and reacting to** platform level events

Aspect-orientation at Platform-level



Aspect-orientation at Platform-level

□ Example implementation



Summary

- It is important to determine and understand platform characteristics
 - Cloud and virtualization based platforms have several unique characteristics
- Devise solutions to application scenarios
 - Leverage platform characteristics