

Machine Learning

Classification Methods

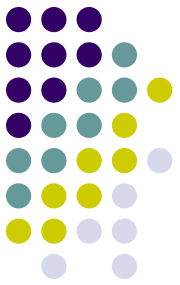
Bayesian Classification, Nearest
Neighbor, Ensemble Methods

Bayesian Classification: Why?



- A statistical classifier: performs *probabilistic prediction*, *i.e.*, predicts class membership probabilities
- Foundation: Based on Bayes' Theorem.
- Performance: A simple Bayesian classifier, *naïve Bayesian classifier*, has comparable performance with decision tree and selected neural network classifiers
- Incremental: Each training example can incrementally increase/decrease the probability that a hypothesis is correct — prior knowledge can be combined with observed data

Bayes' Rule



$$p(h | d) = \frac{P(d | h)P(h)}{P(d)}$$

Understanding Bayes' rule

d = data

h = hypothesis (model)

- rearranging

$$p(h | d)P(d) = P(d | h)P(h)$$

$$P(d, h) = P(d, h)$$

the same joint probability

on both sides

Who is who in Bayes' rule

$P(h)$: prior belief (probability of hypothesis h before seeing any data)

$P(d | h)$: likelihood (probability of the data if the hypothesis h is true)

$P(d) = \sum_h P(d | h)P(h)$: data evidence (marginal probability of the data)

$P(h | d)$: posterior (probability of hypothesis h after having seen the data d)



Example of Bayes Theorem

- Given:
 - A doctor knows that meningitis causes stiff neck 50% of the time
 - Prior probability of any patient having meningitis is $1/50,000$
 - Prior probability of any patient having stiff neck is $1/20$
- If a patient has stiff neck, what's the probability he/she has meningitis?

$$P(M | S) = \frac{P(S | M)P(M)}{P(S)} = \frac{0.5 \times 1/50000}{1/20} = 0.0002$$

Choosing Hypotheses




- *Maximum Likelihood* hypothesis:

$$h_{ML} = \arg \max_{h \in H} P(d | h)$$

- Generally we want the most probable hypothesis given training data. This is the *maximum a posteriori* hypothesis:
 - Useful observation: it does not depend on the denominator $P(d)$


$$h_{MAP} = \arg \max_{h \in H} P(h | d)$$

Bayesian Classifiers



- Consider each attribute and class label as random variables
- Given a record with attributes (A_1, A_2, \dots, A_n)
 - Goal is to predict class C
 - Specifically, we want to find the value of C that maximizes $P(C | A_1, A_2, \dots, A_n)$
- Can we estimate $P(C | A_1, A_2, \dots, A_n)$ directly from data?

Bayesian Classifiers



- Approach:
 - compute the posterior probability $P(C \mid A_1, A_2, \dots, A_n)$ for all values of C using the Bayes theorem

$$P(C \mid A_1 A_2 \dots A_n) = \frac{P(A_1 A_2 \dots A_n \mid C) P(C)}{P(A_1 A_2 \dots A_n)}$$

- Choose value of C that maximizes $P(C \mid A_1, A_2, \dots, A_n)$
- Equivalent to choosing value of C that maximizes $P(A_1, A_2, \dots, A_n \mid C) P(C)$
- How to estimate $P(A_1, A_2, \dots, A_n \mid C)$?



Naïve Bayes Classifier

- Assume independence among attributes A_i when class is given:
 - $P(A_1, A_2, \dots, A_n | C) = P(A_1 | C_j) P(A_2 | C_j) \dots P(A_n | C_j)$
 - Can estimate $P(A_i | C_j)$ for all A_i and C_j .
 - This is a simplifying assumption which may be violated in reality
- The Bayesian classifier that uses the Naïve Bayes assumption and computes the MAP hypothesis is called Naïve Bayes classifier

$$c_{Naive\ Bayes} = \arg \max_c P(c)P(\mathbf{x} | c) = \arg \max_c P(c) \prod_i P(a_i | c)$$

How to Estimate Probabilities from Data?



Tid	Refund	Marital Status	Taxable Income	Evade
1	Yes	Single	125K	No
2	No	Married	100K	No
3	No	Single	70K	No
4	Yes	Married	120K	No
5	No	Divorced	95K	Yes
6	No	Married	60K	No
7	Yes	Divorced	220K	No
8	No	Single	85K	Yes
9	No	Married	75K	No
10	No	Single	90K	Yes

- Class: $P(C) = N_c / N$

- e.g., $P(\text{No}) = 7/10$,
 $P(\text{Yes}) = 3/10$

- For discrete attributes:

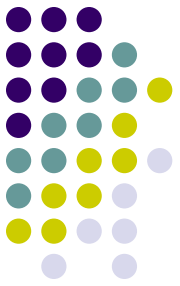
$$P(A_i | C_k) = |A_{ik}| / N_{C_k}$$

- where $|A_{ik}|$ is number of instances having attribute A_i and belongs to class C_k
 - Examples:

$$P(\text{Status}=\text{Married} | \text{No}) = 4/7$$


$$P(\text{Refund}=\text{Yes} | \text{Yes})=0$$

How to Estimate Probabilities from Data?



- For continuous attributes:
 - **Discretize** the range into bins
 - one ordinal attribute per bin
 - violates independence assumption
 - **Two-way split:** $(A < v)$ or $(A > v)$
 - choose only one of the two splits as new attribute
 - **Probability density estimation:**
 - Assume attribute follows a normal distribution
 - Use data to estimate parameters of distribution (e.g., mean and standard deviation)
 - Once probability distribution is known, can use it to estimate the conditional probability $P(A_i | c)$

How to Estimate Probabilities from Data?



Tid	Refund	Marital Status	Taxable Income	Evade
1	Yes	Single	125K	No
2	No	Married	100K	No
3	No	Single	70K	No
4	Yes	Married	120K	No
5	No	Divorced	95K	Yes
6	No	Married	60K	No
7	Yes	Divorced	220K	No
8	No	Single	85K	Yes
9	No	Married	75K	No
10	No	Single	90K	Yes

- Normal distribution:

$$P(A_i | c_j) = \frac{1}{\sqrt{2\pi\sigma_{ij}^2}} e^{-\frac{(A_i - \mu_{ij})^2}{2\sigma_{ij}^2}}$$

- One for each (A_i, c_i) pair
- For (Income, Class=No):
 - If Class=No
 - sample mean = 110
 - sample variance = 2975

$$P(\text{Income} = 120 | \text{No}) = \frac{1}{\sqrt{2\pi(54.54)}} e^{-\frac{(120-110)^2}{2(2975)}} = 0.0072$$

Naïve Bayesian Classifier: Training Dataset



Class:

C1:buys_computer = 'yes'

C2:buys_computer = 'no'

New Data:

X = (age <=30,

Income = medium,

Student = yes

Credit_rating = Fair)

age	income	student	credit_rating	com
<=30	high	no	fair	no
<=30	high	no	excellent	no
31...40	high	no	fair	yes
>40	medium	no	fair	yes
>40	low	yes	fair	yes
>40	low	yes	excellent	no
31...40	low	yes	excellent	yes
<=30	medium	no	fair	no
<=30	low	yes	fair	yes
>40	medium	yes	fair	yes
<=30	medium	yes	excellent	yes
31...40	medium	no	excellent	yes
31...40	high	yes	fair	yes
>40	medium	no	excellent	no

Naïve Bayesian Classifier: An Example



Given X (age=youth, income=medium, student=yes, credit=fair)

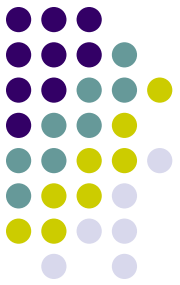
Maximize $P(X|C_i)P(C_i)$, for $i=1,2$

First step: Compute $P(C)$ The prior probability of each class can be computed based on the training tuples:

$$P(\text{buys_computer=yes}) = 9/14 = 0.643$$

$$P(\text{buys_computer=no}) = 5/14 = 0.357$$

Naïve Bayesian Classifier: An Example



Given X (age=youth, income=medium, student=yes, credit=fair)

Maximize $P(X|C_i)P(C_i)$, for $i=1,2$

Second step: compute $P(X|C_i)$

$$\begin{aligned} P(X|\text{buys_computer=yes}) &= P(\text{age=youth} | \text{buys_computer=yes}) \times \\ &\quad P(\text{income=medium} | \text{buys_computer=yes}) \times \\ &\quad P(\text{student=yes} | \text{buys_computer=yes}) \times \\ &\quad P(\text{credit_rating=fair} | \text{buys_computer=yes}) \\ &= 0.044 \end{aligned}$$


$$P(\text{age=youth} | \text{buys_computer=yes}) = 0.222$$

$$P(\text{income=medium} | \text{buys_computer=yes}) = 0.444$$

$$P(\text{student=yes} | \text{buys_computer=yes}) = 6/9 = 0.667$$

$$P(\text{credit_rating=fair} | \text{buys_computer=yes}) = 6/9 = 0.667$$

Naïve Bayesian Classifier: An Example



Given X (age=youth, income=medium, student=yes, credit=fair)

Maximize $P(X|C_i)P(C_i)$, for $i=1,2$

Second step: compute $P(X|C_i)$

$P(X|\text{buys_computer=no}) = P(\text{age=youth} | \text{buys_computer=no}) \times$

$P(\text{income=medium} | \text{buys_computer=no}) \times$

$P(\text{student=yes} | \text{buys_computer=no}) \times$

$P(\text{credit_rating=fair} | \text{buys_computer=no})$

$= 0.019$


$P(\text{age=youth} | \text{buys_computer=no}) = 3/5 = 0.666$

$P(\text{income=medium} | \text{buys_computer=no}) = 2/5 = 0.400$

$P(\text{student=yes} | \text{buys_computer=no}) = 1/5 = 0.200$

$P(\text{credit_rating=fair} | \text{buys_computer=no}) = 2/5 = 0.400$

Naïve Bayesian Classifier: An Example



Given X (age=youth, income=medium, student=yes, credit=fair)

Maximize $P(X|C_i)P(C_i)$, for $i=1,2$

We have computed in the first and second steps:

$$P(\text{buys_computer}=\text{yes})=9/14=0.643$$

$$P(\text{buys_computer}=\text{no})=5/14=0.357$$

$$P(X|\text{buys_computer}=\text{yes})= 0.044$$

$$P(X|\text{buys_computer}=\text{no})= 0.019$$

Third step: compute $P(X|C_i)P(C_i)$ for each class

$$P(X|\text{buys_computer}=\text{yes})P(\text{buys_computer}=\text{yes})=0.044 \times 0.643=0.028$$

$$P(X|\text{buys_computer}=\text{no})P(\text{buys_computer}=\text{no})=0.019 \times 0.357=0.007$$

The naïve Bayesian Classifier predicts **X belongs to class (“buys_computer = yes”)**

Example



Training set :
(Öğrenme Kümesi)

<i>Tid</i>	Refund	Marital Status	Taxable Income	Evade
1	Yes	Single	125K	No
2	No	Married	100K	No
3	No	Single	70K	No
4	Yes	Married	120K	No
5	No	Divorced	95K	Yes
6	No	Married	60K	No
7	Yes	Divorced	220K	No
8	No	Single	85K	Yes
9	No	Married	75K	No
10	No	Single	90K	Yes

Given a Test Record:

$X = (\text{Refund} = \text{No}, \text{Married}, \text{Income} = 120\text{K})$

k

Example of Naïve Bayes Classifier



Given a Test Record:

$X = (\text{Refund} = \text{No}, \text{Married}, \text{Income} = 120\text{K})$

naive Bayes Classifier:

$P(\text{Refund}=\text{Yes}|\text{No}) = 3/7$
 $P(\text{Refund}=\text{No}|\text{No}) = 4/7$
 $P(\text{Refund}=\text{Yes}|\text{Yes}) = 0$
 $P(\text{Refund}=\text{No}|\text{Yes}) = 1$
 $P(\text{Marital Status}=\text{Single}|\text{No}) = 2/7$
 $P(\text{Marital Status}=\text{Divorced}|\text{No}) = 1/7$
 $P(\text{Marital Status}=\text{Married}|\text{No}) = 4/7$
 $P(\text{Marital Status}=\text{Single}|\text{Yes}) = 2/7$
 $P(\text{Marital Status}=\text{Divorced}|\text{Yes}) = 1/7$
 $P(\text{Marital Status}=\text{Married}|\text{Yes}) = 0$

For taxable income:

If class=No: sample mean=110
 sample variance=2975
If class=Yes: sample mean=90
 sample variance=25

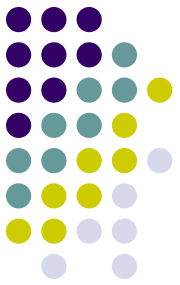
- $P(X|\text{Class}=\text{No}) = P(\text{Refund}=\text{No}|\text{Class}=\text{No})$
 $\times P(\text{Married}|\text{Class}=\text{No})$
 $\times P(\text{Income}=120\text{K}|\text{Class}=\text{No})$
 $= 4/7 \times 4/7 \times 0.0072 = 0.0024$
- $P(X|\text{Class}=\text{Yes}) = P(\text{Refund}=\text{No}|\text{Class}=\text{Yes})$
 $\times P(\text{Married}|\text{Class}=\text{Yes})$
 $\times P(\text{Income}=120\text{K}|\text{Class}=\text{Yes})$
 $= 1 \times 0 \times 1.2 \times 10^{-9} = 0$

Since $P(X|\text{No})P(\text{No}) > P(X|\text{Yes})P(\text{Yes})$

Therefore $P(\text{No}|X) > P(\text{Yes}|X)$

$\Rightarrow \text{Class} = \text{No}$

Avoiding the 0-Probability Problem



- If one of the conditional probability is zero, then the entire expression becomes zero
- Probability estimation:

$$\text{Original : } P(A_i | C) = \frac{N_{ic}}{N_c}$$

$$\text{Laplace : } P(A_i | C) = \frac{N_{ic} + 1}{N_c + c}$$

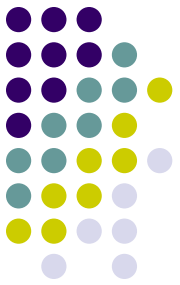
$$\text{m - estimate : } P(A_i | C) = \frac{N_{ic} + mp}{N_c + m}$$

c: number of classes

p: prior probability

m: parameter

Naïve Bayes (Summary)



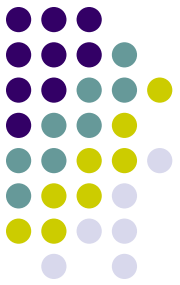
- Advantage
 - Robust to isolated noise points
 - Handle missing values by ignoring the instance during probability estimate calculations
 - Robust to irrelevant attributes
- Disadvantage
 - Assumption: class conditional independence, which may cause loss of accuracy
 - Independence assumption may not hold for some attribute. Practically, dependencies exist among variables
 - Use other techniques such as Bayesian Belief Networks (BBN)

Remember



- Bayes' rule can be turned into a classifier
- Maximum A Posteriori (MAP) hypothesis estimation incorporates prior knowledge; Max Likelihood (ML) doesn't
- Naive Bayes Classifier is a simple but effective Bayesian classifier for vector data (i.e. data with several attributes) that assumes that attributes are independent given the class.
- Bayesian classification is a generative approach to classification

Classification Paradigms




- In fact, we can categorize three fundamental approaches to classification:
- **Generative models:** Model $p(x|C_k)$ and $P(C_k)$ separately and use the Bayes theorem to find the posterior probabilities $P(C_k|x)$
 - E.g. Naive Bayes, Gaussian Mixture Models, Hidden Markov Models,...
- **Discriminative models:**
 - Determine $P(C_k|x)$ directly and use in decision
 - E.g. Linear discriminant analysis, SVMs, NNs,...
- Find a **discriminant function** f that maps x onto a class label directly without calculating probabilities




Bayesian Belief Networks

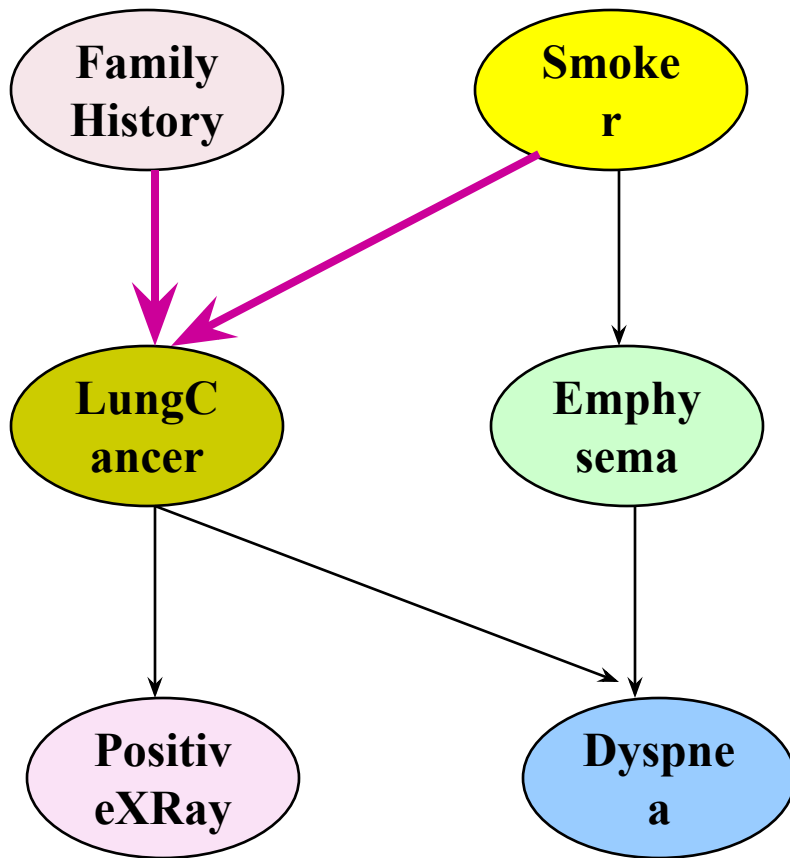
- Bayesian belief network allows a *subset* of the variables to be conditionally independent
- A graphical model of causal relationships (*neden sonuç ilişkilerini simgeleyen bir çizge tabanlı model*)
 - Represents dependency among the variables
 - Gives a specification of joint probability distribution



- ☐ **Nodes:** random variables
- ☐ **Links:** dependency
- ☐ X and Y are the parents of Z, and Y is the parent of P
- ☐ No dependency between Z and P
- ☐ Has no loops or cycles



Bayesian Belief Network: An Example



The **conditional probability table (CPT)** for variable LungCancer:

	(FH, S)	(FH, ~S)	(~FH, S)	(~FH, ~S)
LC	0.8	0.5	0.7	0.1
~LC	0.2	0.5	0.3	0.9

CPT shows the conditional probability for each possible combination of its parents

Derivation of the probability of a particular combination of values of **X**, from CPT:

Bayesian Belief Networks

$$P(x_1, \dots, x_n) = \prod_{i=1}^n P(x_i \mid \text{Parents}(Y_i))$$

Training Bayesian Networks



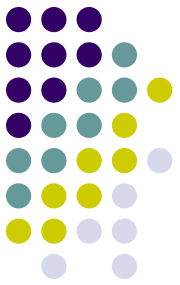
- Several scenarios:
 - Given both the network structure and all variables observable: *learn only the CPTs*
 - Network structure known, some hidden variables: *gradient descent* (greedy hill-climbing) method, analogous to neural network learning
 - Network structure unknown, all variables observable: search through the model space to *reconstruct network topology*
 - Unknown structure, all hidden variables: No good algorithms known for this purpose
- Ref. D. Heckerman: Bayesian networks for data mining

Lazy Learners



- The classification algorithms presented before are **eager learners**
 - Construct a model before receiving new tuples to classify
 - Learned models are ready and eager to classify previously unseen tuples
- **Lazy learners**
 - The learner waits till the last minute before doing any model construction
 - In order to classify a given test tuple
 - Store training tuples
 - Wait for test tuples
 - Perform generalization based on similarity between test and the stored training tuples

Lazy vs Eager




Eager Learners	Lazy Learners
<ul style="list-style-type: none">• Do lot of work on training data	<ul style="list-style-type: none">• Do less work on training data
<ul style="list-style-type: none">• Do less work when test tuples are presented	<ul style="list-style-type: none">• Do more work when test tuples are presented

Basic k-Nearest Neighbor Classification



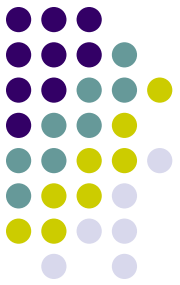
- Given training data $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_N, y_N)$
- Define a distance metric between points in input space $D(\mathbf{x}_1, \mathbf{x}_i)$
 - E.g., Euclidean distance, Weighted Euclidean, Mahalanobis distance, TFIDF, etc.
- Training method:
 - Save the training examples
- At prediction time:
 - Find the k training examples $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_k, y_k)$ that are closest to the test example \mathbf{x} given the distance $D(\mathbf{x}_1, \mathbf{x}_i)$
 - Predict the most frequent class among those y_i 's.

Nearest-Neighbor Classifiers



- Requires three things
 - The set of stored records
 - Distance Metric to compute distance between records
 - The value of k , the number of nearest neighbors to retrieve
- To classify an unknown record:
 - Compute distance to other training records
 - Identify k nearest neighbors
 - Use class labels of nearest neighbors to determine the class label of unknown record (e.g., by taking majority vote)

K-Nearest Neighbor Model




- **Classification:**

\hat{y} = most common class in set $\{y_1, \dots, y_K\}$

- **Regression:**

$$\hat{y} = \frac{1}{K} \sum_{k=1}^K y_k$$

K-Nearest Neighbor Model: Weighted by Distance



- **Classification:**



\hat{y} = most common class in weighted set

$$\{D(\mathbf{x}, \mathbf{x}_1)y_1, \dots, D(\mathbf{x}, \mathbf{x}_K)y_K\}$$


- **Regression:**

$$\hat{y} = \frac{\sum_{k=1}^K D(x, x_k) y_k}{\sum_{k=1}^K D(x, x_k)}$$


Definition of Nearest Neighbor



(a) 1-nearest neighbor



(b) 2-nearest neighbor




(c) 3-nearest neighbor

K-nearest neighbors of a record x are data points that have the k smallest distance to x


The decision boundary implemented by 3NN



The boundary is always the perpendicular bisector of the line between two points (Voronoi tessellation)



(a)




(b)

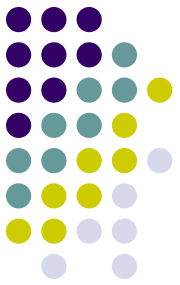
Nearest Neighbor Classification...



- Choosing the value of k :
 - If k is too small, sensitive to noise points
 - If k is too large, neighborhood may include points from other classes



Determining the value of k



- In typical applications k is in units or tens rather than in hundreds or thousands
- Higher values of k provide smoothing that reduces the risk of overfitting due to noise in the training data
- Value of k can be chosen based on error rate measures
- We should also avoid over-smoothing by choosing $k=n$, where n is the total number of tuples in the training data set



Determining the value of k

- Given training examples $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_N, y_N)$
- Use N fold cross validation
 - Search over $K = (1, 2, 3, \dots, Kmax)$. Choose search size $Kmax$ based on compute constraints
 - Calculated the average error for each K:
 - Calculate predicted class \hat{y}_i for each training point (\mathbf{x}_i, y_i) , $i = 1, \dots, N$
(using all other points to build the model)
 - Average over all training examples
- Pick K to minimize the cross validation error

Example



RID	Income(\$000's)	lot Size (000's sq.ft)	class: Owners =1 Non-Owners=2
1	60	18.4	1
2	85.5	16.8	1
3	64.8	21.6	1
4	61.5	20.8	1
5	87	23.6	1
6	110.1	19.2	1
7	108	17.6	1
8	82.8	22.4	1
9	69	20	1
10	93	20.8	1
11	51	22	1
12	81	20	2
13	75	19.6	2
14	52.8	20.8	2
15	64.8	17.2	2
16	43.2	20.4	2
17	84	17.6	2
18	49.2	17.6	2
19	59.4	16	2
20	66	18.4	2
21	47.4	16.4	2
22	33	18.8	2
23	51	14	2
24	63	14.8	2



We randomly divide the data into

18 training cases

6 test cases:

tuples 6,7,12,14,19, 20

Use training cases to classify test cases and compute error rates

Choosing k




- ▶ If we choose $k=1$ we will classify in a way that is very sensitive to the local characteristics of our data
- ▶ If we choose a **large value of k** we average over a large number of data points and average out the variability due to the noise associated with data points
- ▶ If we choose $k=18$ we would simply predict the most frequent class in the data set in all cases
 - Very stable but completely ignores the information in the independent variables

k	1	3	5	7	9	11	13	18
Misclassification error %	33	33	33	33	33	17	17	50


- We would choose $k=11$ (or possibly 13) in this case

Nearest neighbor Classification...



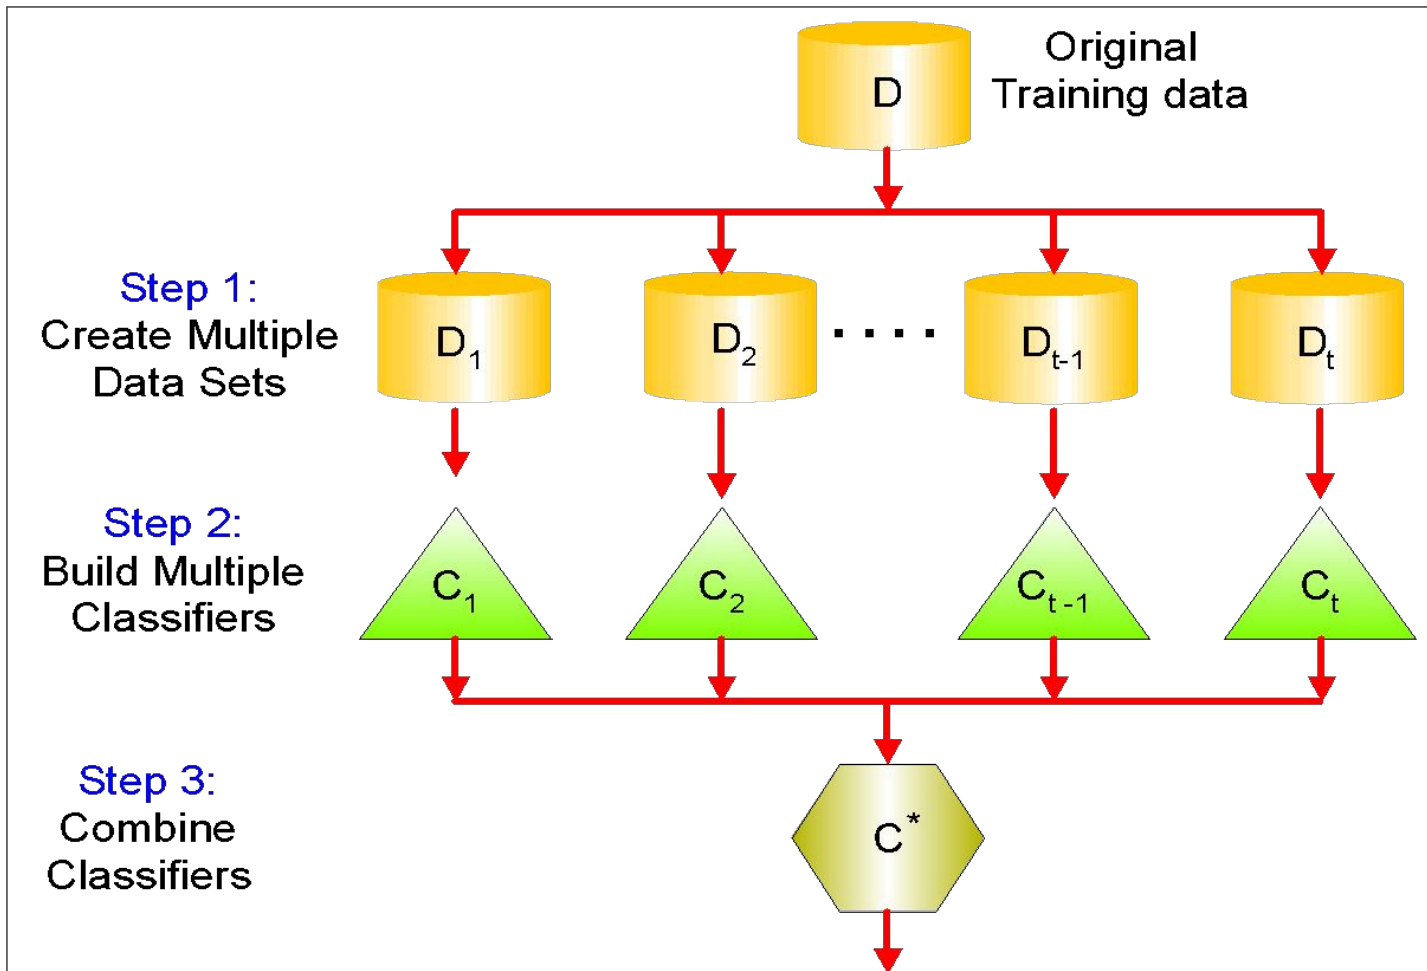
- k-NN classifiers are lazy learners
 - It does not build models explicitly
 - Unlike eager learners such as decision tree induction and rule-based systems
- Adv: No training time
- Disadv:
 - Testing time can be long, classifying unknown records are relatively expensive
 - Curse of Dimensionality : Can be easily fooled in high dimensional spaces
 - Dimensionality reduction techniques are often used

Ensemble Methods



- One of the eager methods => builds model over the training set
- Construct a set of classifiers from the training data
- Predict class label of previously unseen records by aggregating predictions made by multiple classifiers

General Idea



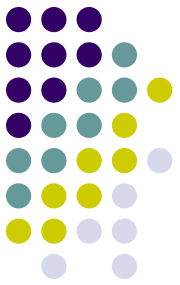
Why does it work?



- Suppose there are 25 base classifiers
 - Each classifier has error rate, $\varepsilon = 0.35$
 - Assume classifiers are independent
 - Probability that the ensemble classifier makes a wrong prediction:

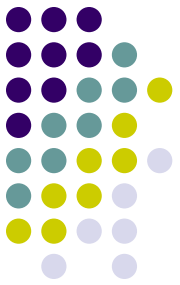
$$\sum_{i=1}^{25} \binom{25}{i} \varepsilon^i (1 - \varepsilon)^{25-i} = 0.06$$

Examples of Ensemble Methods



- How to generate an ensemble of classifiers?
 - Bagging
 - Boosting
 - Random Forests

Bagging: Bootstrap AGGREGatING



- Bootstrap: data resampling
 - Generate multiple training sets
 - Resample the original training data
 - With replacement
 - Data sets have different “specious” patterns
- Sampling with replacement
 - Each sample has probability $(1 - 1/n)^n$ of being selected

Original Data	1	2	3	4	5	6	7	8	9	10
Bagging (Round 1)	7	8	10	8	2	5	10	10	5	9
Bagging (Round 2)	1	4	9	1	2	3	2	7	3	2
Bagging (Round 3)	1	8	5	10	5	5	9	6	3	7

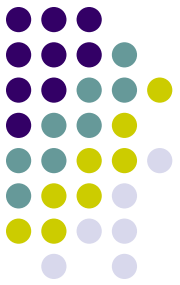
- Build classifier on each bootstrap sample
 - Specious patterns will not correlate
- Underlying true pattern will be common to many
- Combine the classifiers: Label new test examples by a majority vote among classifiers

Boosting



- An iterative procedure to adaptively change distribution of training data by focusing more on previously misclassified records
 - Initially, all N records are assigned equal weights
 - Unlike bagging, weights may change at the end of boosting round
- The final classifier is the weighted combination of the weak classifiers.

Boosting



- Records that are wrongly classified will have their weights increased
- Records that are classified correctly will have their weights decreased

Original Data	1	2	3	4	5	6	7	8	9	10
Boosting (Round 1)	7	3	2	8	7	9	4	10	6	3
Boosting (Round 2)	5	4	9	4	2	5	1	7	4	2
Boosting (Round 3)	4	4	8	10	4	5	4	6	3	4

- Example 4 is hard to classify
- Its weight is increased, therefore it is more likely to be chosen again in subsequent rounds